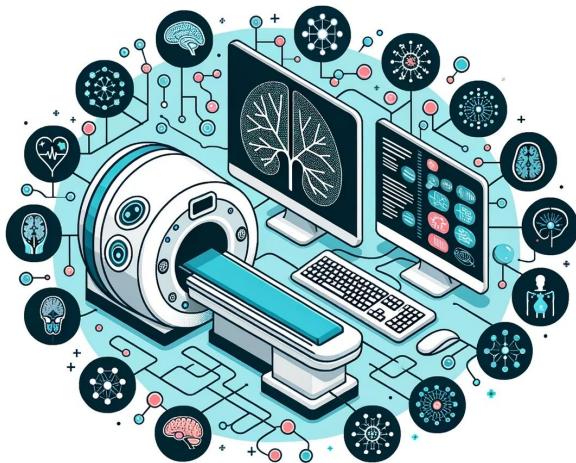


Applied Deep Learning and Generative Models in Healthcare

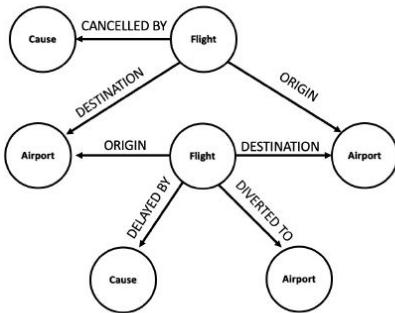


Session 3: Graph neural networks
Date: Jan 25 2025

Instructor: Mahmoud E. Khani, Ph.D.

Many types of data are graphs

- Graphs are a general language for describing and analyzing entities with relations/interactions

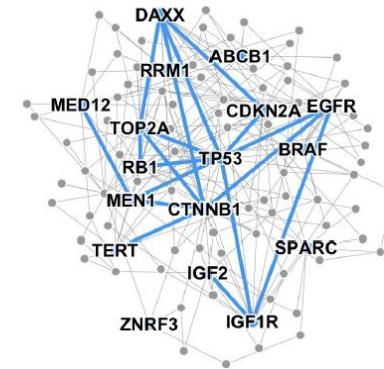


Event Graphs



Image credit: [SalientNetworks](#)

Computer Networks



Disease Pathways

Many types of data are graphs

- Graphs are a general language for describing and analyzing entities with relations/interactions

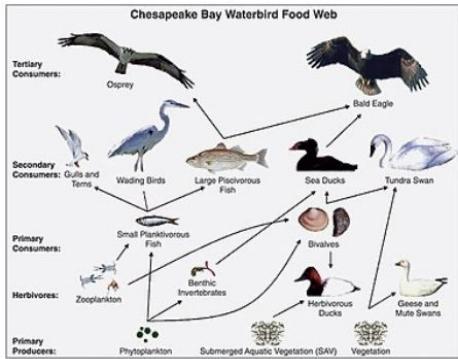


Image credit: [Wikipedia](#)

Food Webs



Image credit: [Pinterest](#)

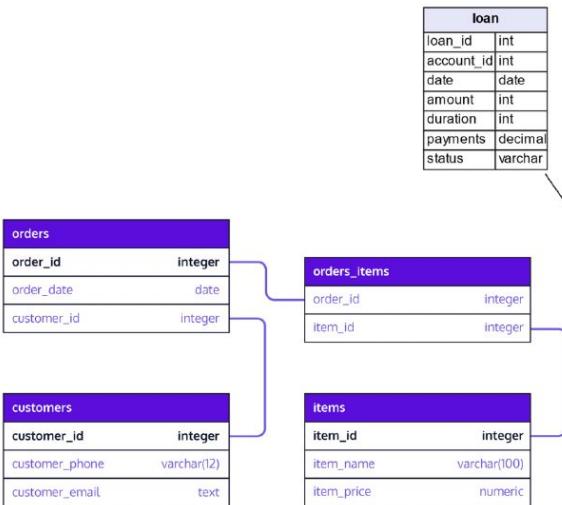
Particle Networks



Image credit: visitlondon.com

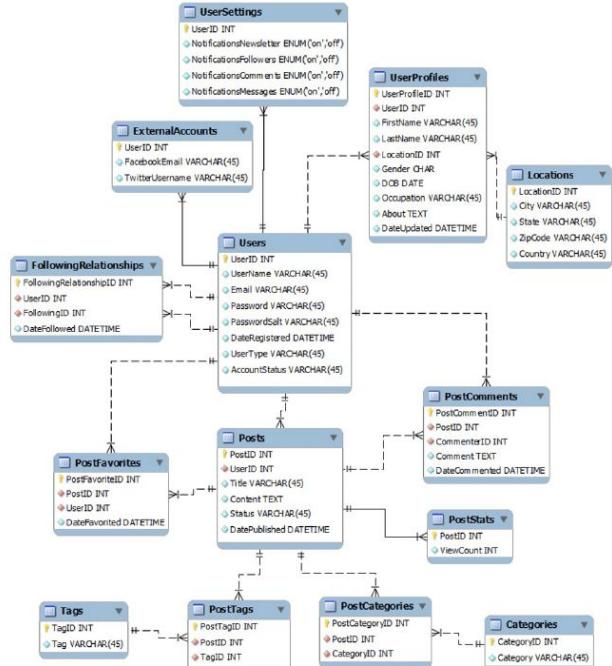
Underground Networks

Databases are Graphs!



Commerce

Finance



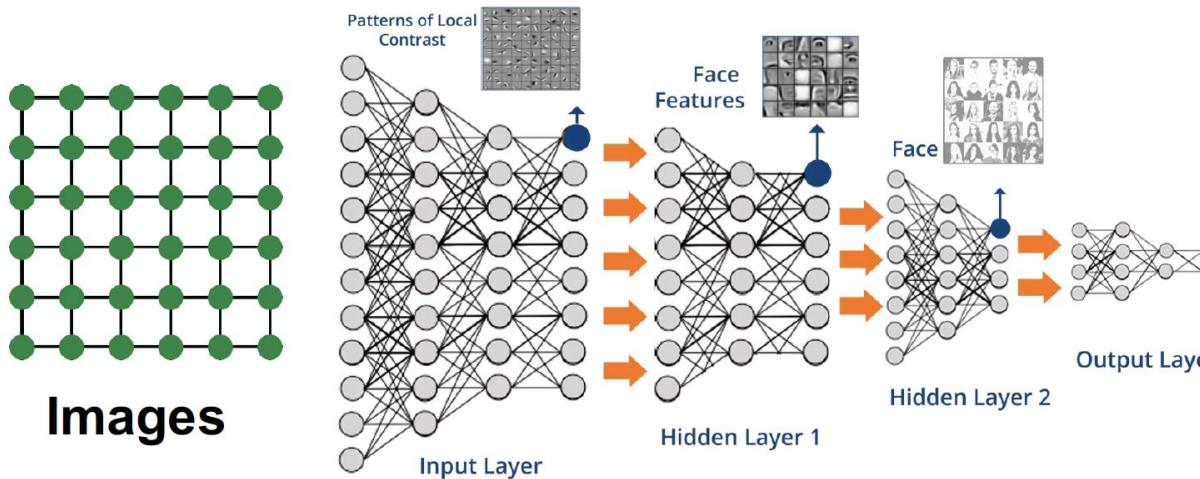
Social Media

Graphs: Machine Learning

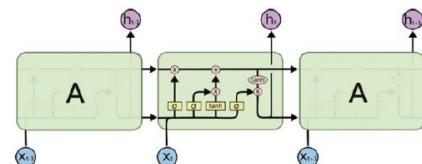
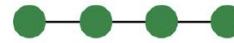
- Complex domains have a rich relational structure, which can be represented as a **relational graph!**
- How do we take advantage of relational structure for better prediction?

Modern ML Toolbox

- Modern deep learning toolbox is designed for simple sequences & grids!



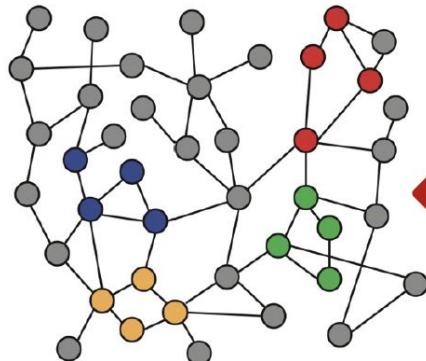
Text/Speech



Why is Graph Deep Learning Hard?

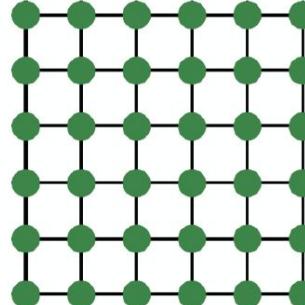
Networks are complex

- Arbitrary size & complex topological structure (i.e., no spatial locality like grids)
- No fixed node ordering or reference point

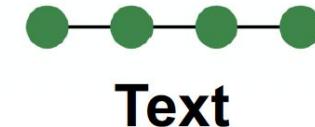


Networks

VS.



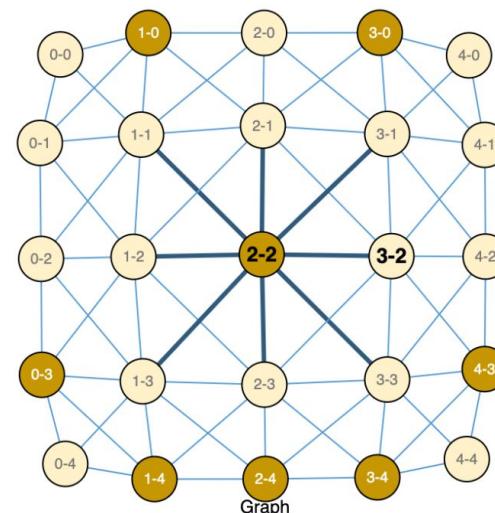
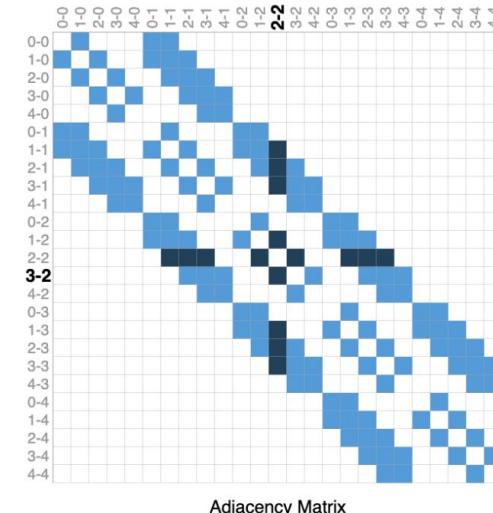
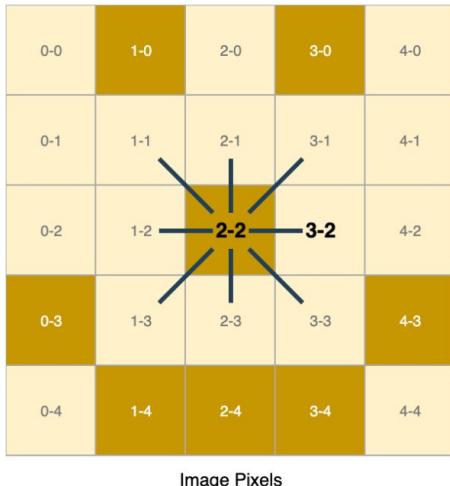
Images



Text

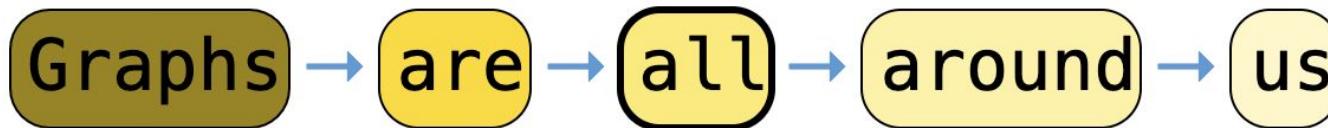
Images as graphs

- Images as rectangular grids with image channels (**244x244x3 floats**)
- images as graphs with regular structure, where each pixel represents a node connected by edges.

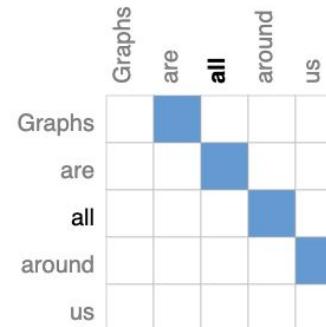


Text as graphs

- We can digitize text by associating indices to each character, word, or token, and representing text as a sequence of these indices.
- This creates a **directed graph!**



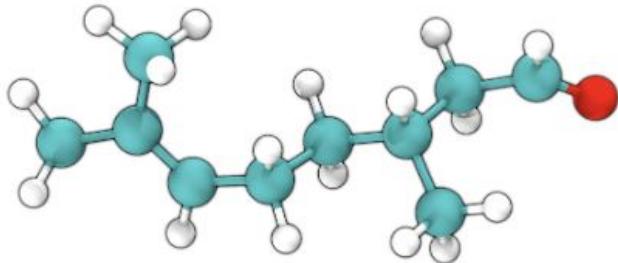
Undirected edge



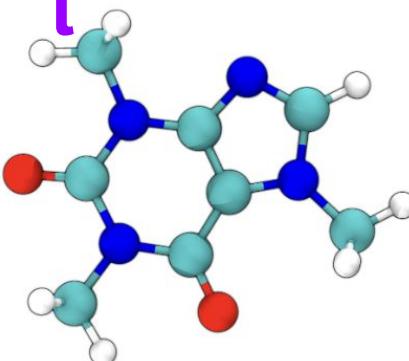
Directed edge



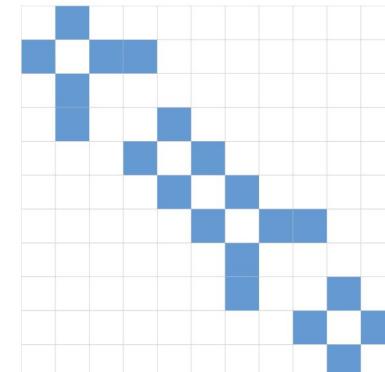
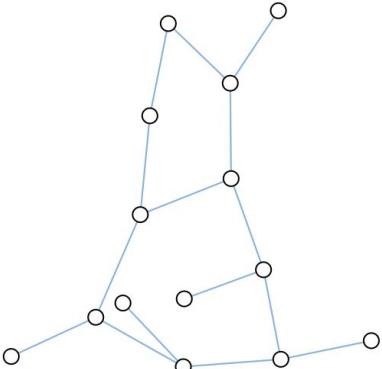
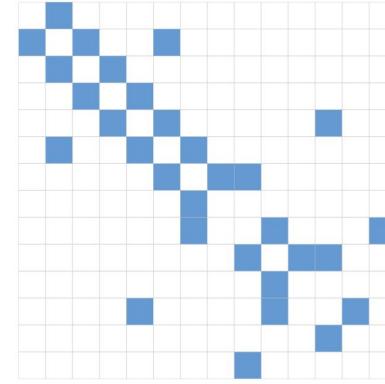
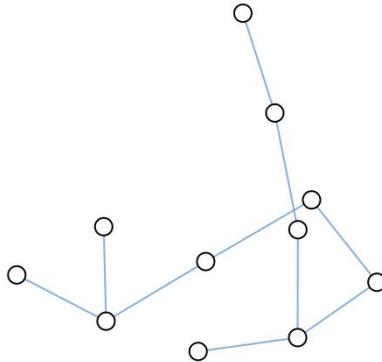
Molecules as graphs



Citronella



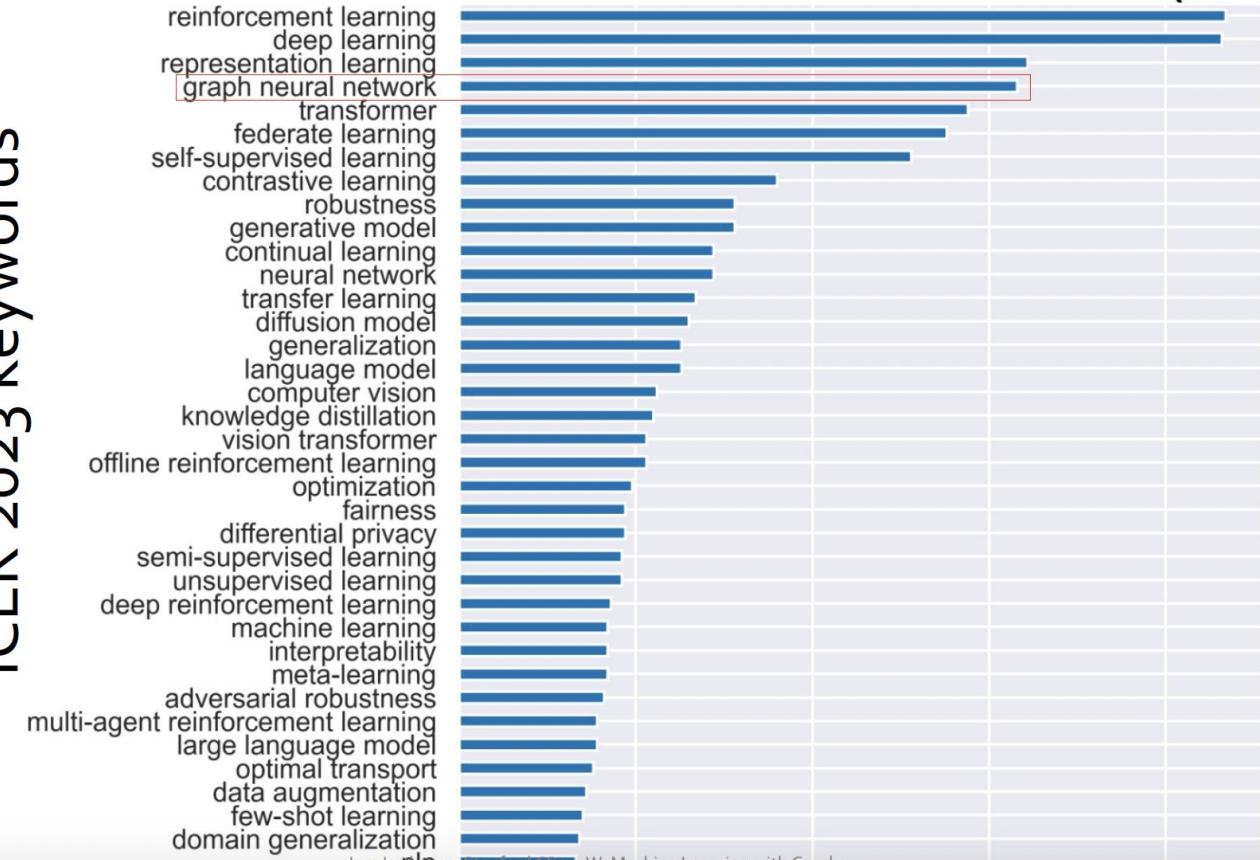
Caffeine



Hot subfiled in ML

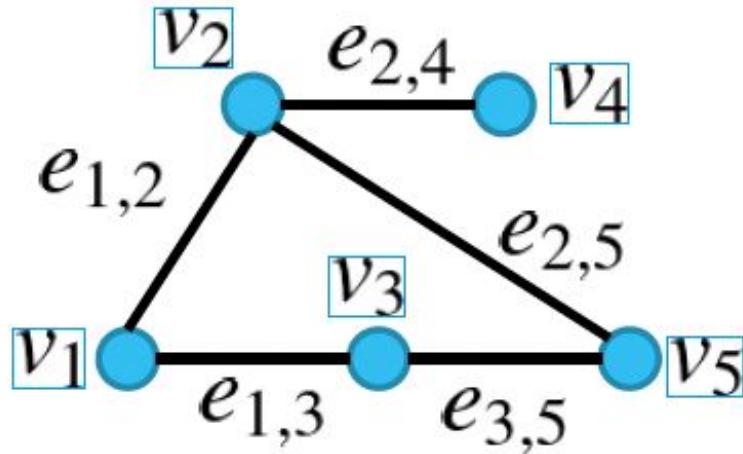
ICLR 2023 keywords

50 MOST APPEARED KEYWORDS (2023)

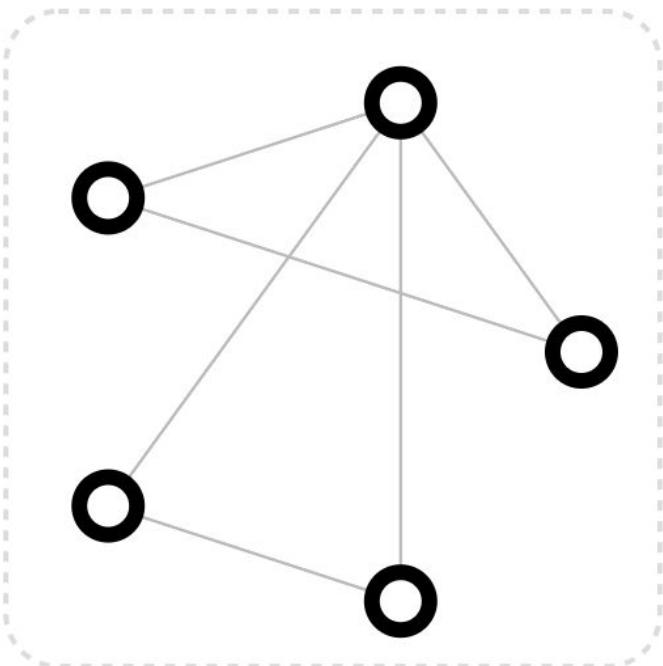


Graph neural networks

- Graph $G = \{v, \varepsilon\}$
- Set of nodes (vertices) $v = \{v_i\}$
- Set of edges $\varepsilon = \{e_{i,j}\}$
 - directional/non-directional
- Nodes and/or edges can have some features
 - Label, numbers, vectors



graph represents the relations (edges) between a collection of entities (nodes)

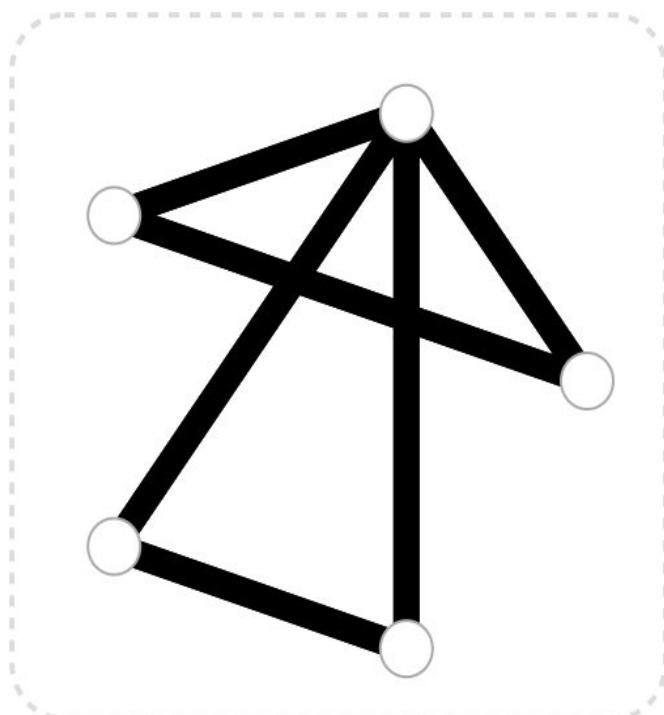


V Vertex (or node) attributes
e.g., node identity, number of neighbors

E Edge (or link) attributes and directions
e.g., edge identity, edge weight

U Global (or master node) attributes
e.g., number of nodes, longest path

graph represents the relations (edges) between a collection of entities (nodes)

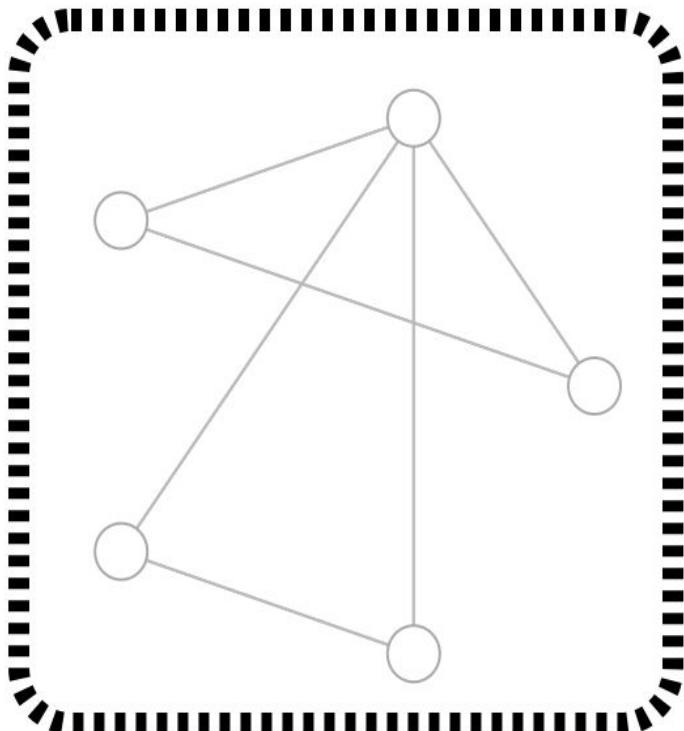


V Vertex (or node) attributes
e.g., node identity, number of neighbors

E Edge (or link) attributes and directions
e.g., edge identity, edge weight

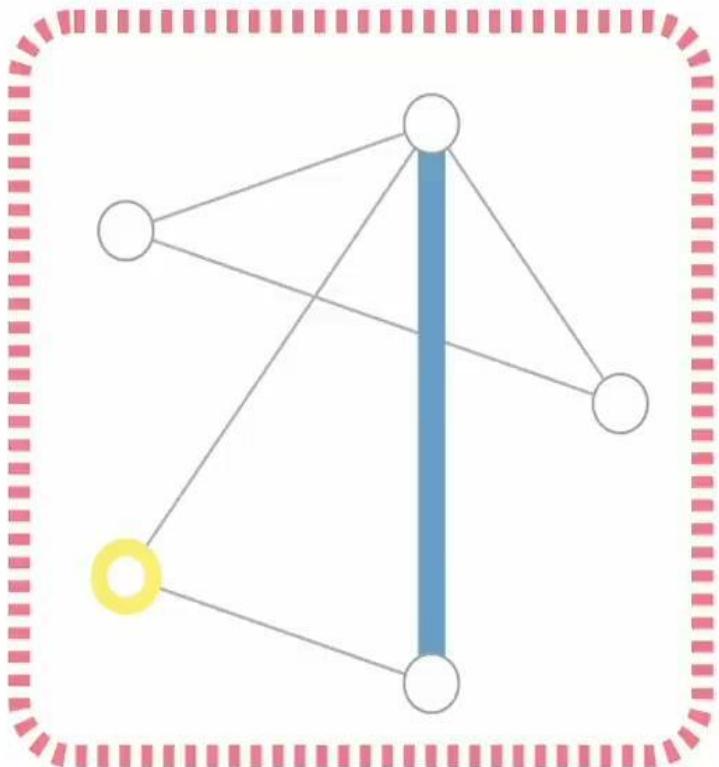
U Global (or master node) attributes
e.g., number of nodes, longest path

graph represents the relations (edges) between a collection of entities (nodes)



- V** Vertex (or node) attributes
 - e.g., node identity, number of neighbors
- E** Edge (or link) attributes and directions
 - e.g., edge identity, edge weight
- U** Global (or master node) attributes
 - e.g., number of nodes, longest path

we can store information in each of pieces of the graph



Vertex (or node) embedding



Edge (or link) attributes and embedding



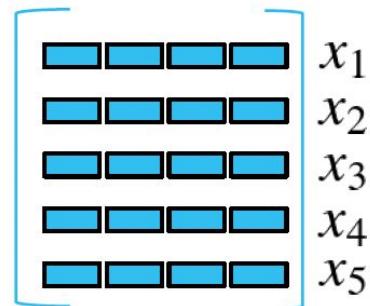
Global (or master node) embedding



Matrix representation of attributed graphs

Feature Matrix $X = (x_{i,d}) = (x_1 \ x_2 \ \dots)^T$

$X \in \mathbb{R}^{N \times D}$ N nodes, D dim. features



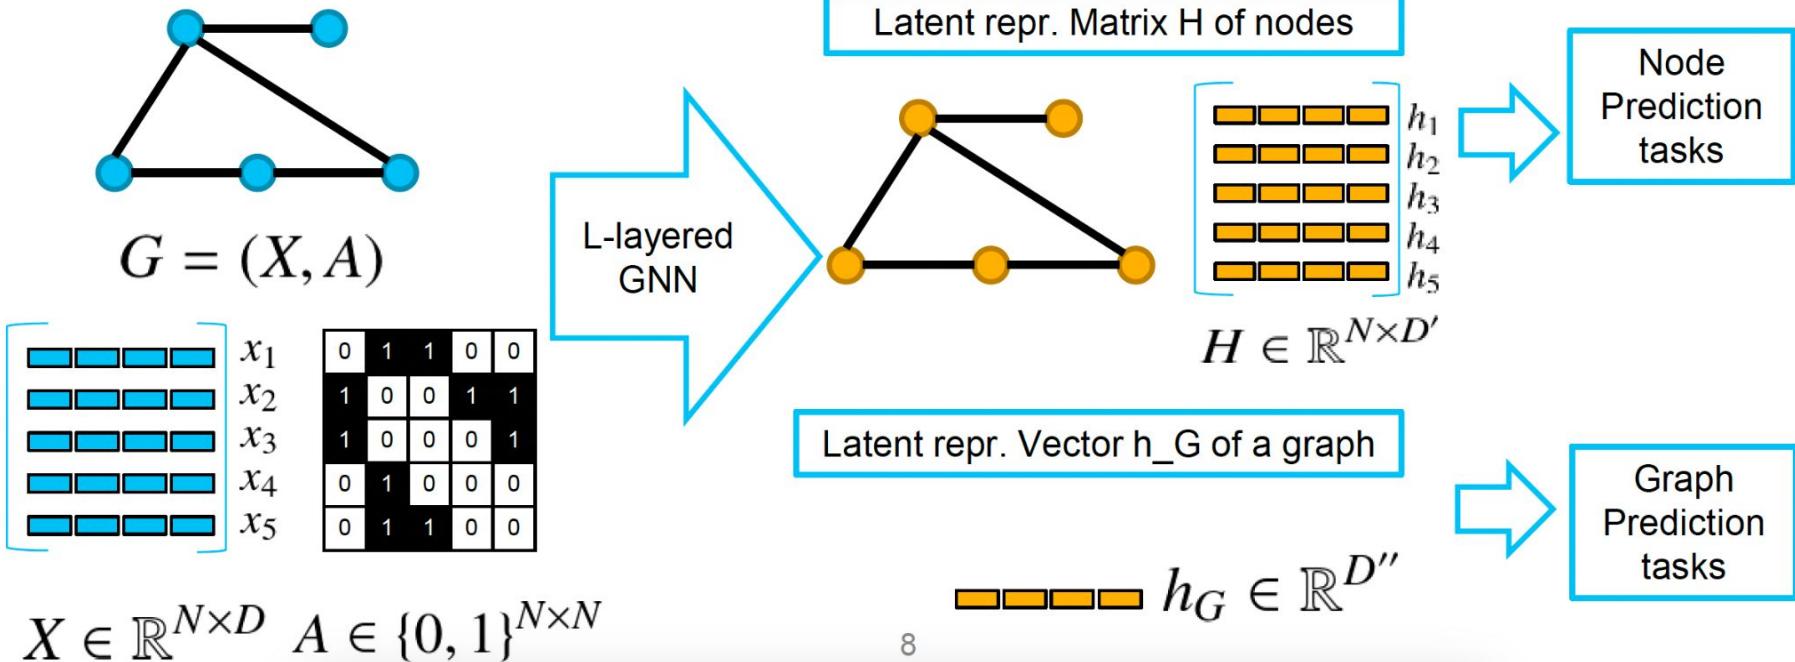
Adjacency Matrix $A = (a_{i,j})_{i,j=1,2,\dots}$

$A \in \{0, 1\}^{N \times N}$ Edge existence between
N X N node pairs

Symmetric A <-- non-directional edges

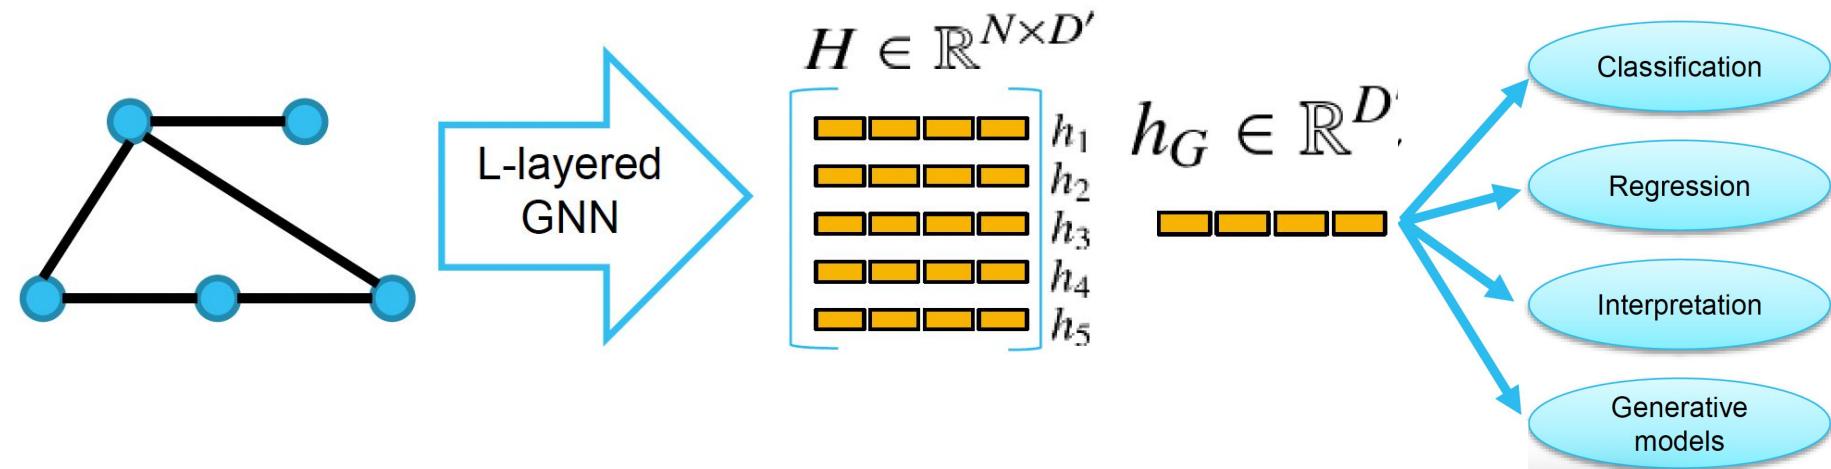
0	1	1	0	0
1	0	0	1	1
1	0	0	0	1
0	1	0	0	0
0	1	1	0	0

GNN functionality: Compute latent representation



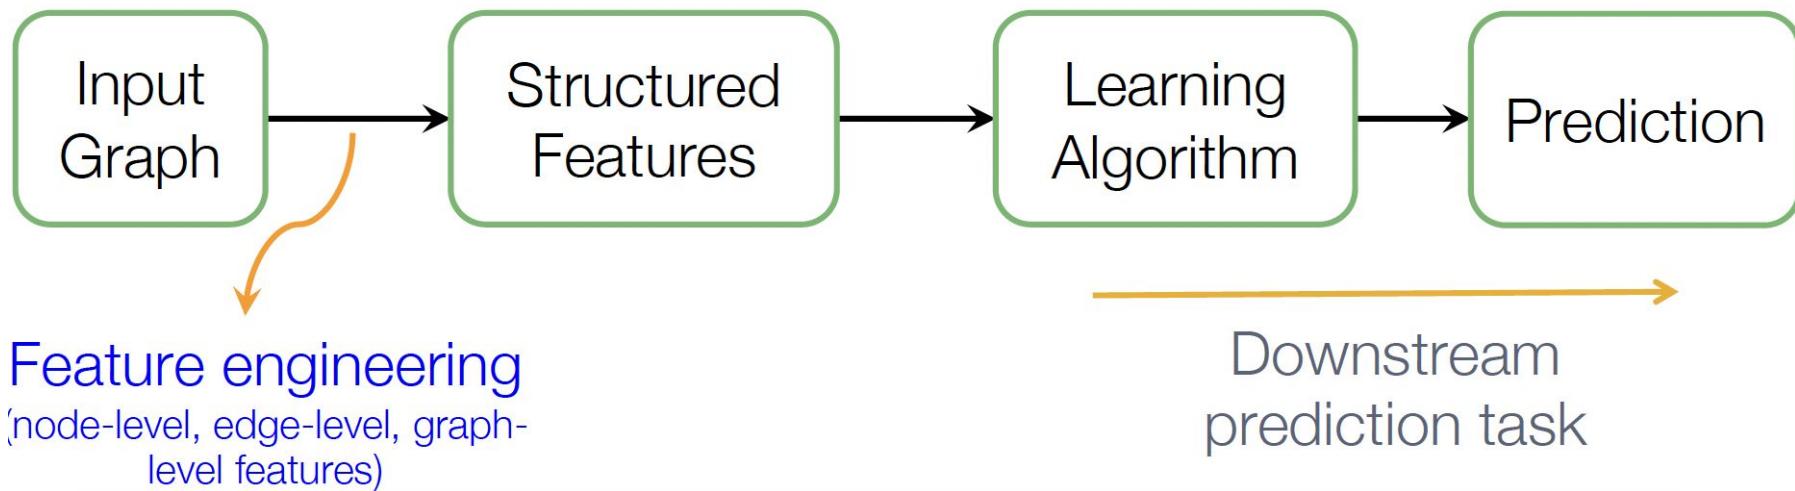
GNN functionality: Compute latent representation

- H is informative and easy-to-handle “data” for machine learning (ML) models/tools.



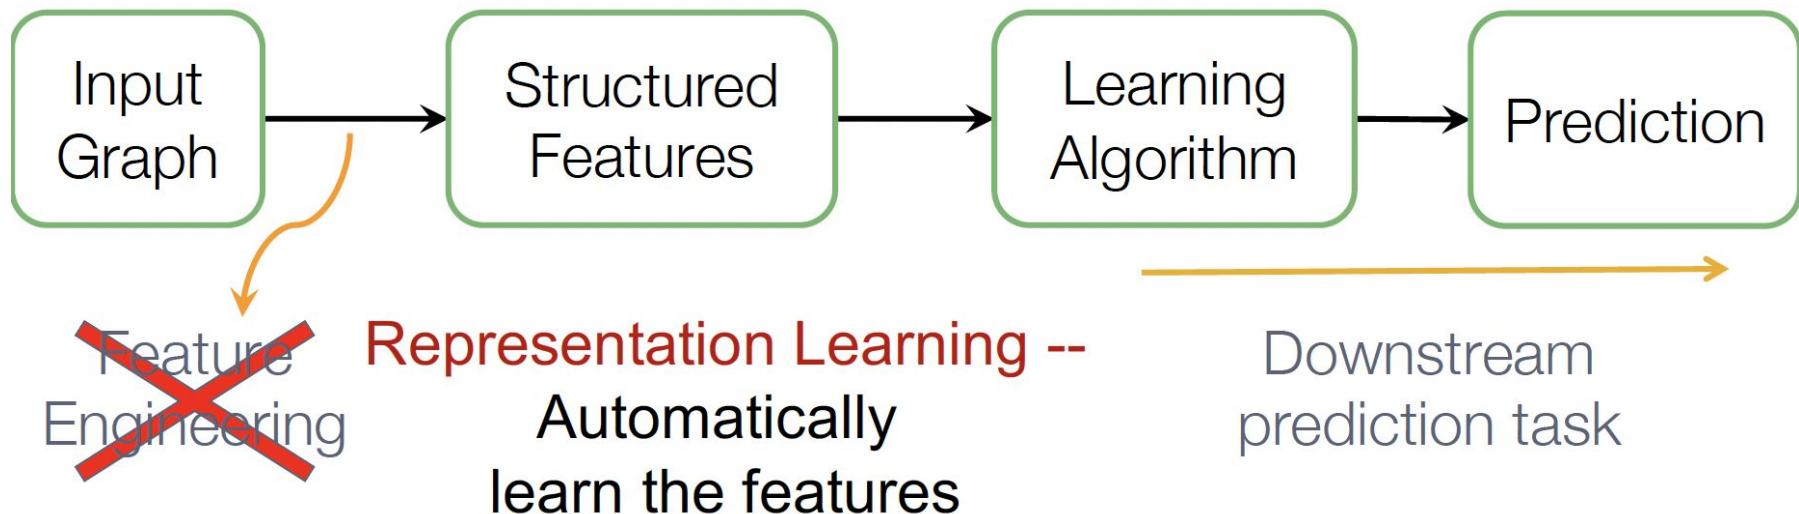
Traditional ML

- Given an input graph, extract node, link and graph-level features, then learn a model (SVM, neural network, etc.) that maps features to labels.



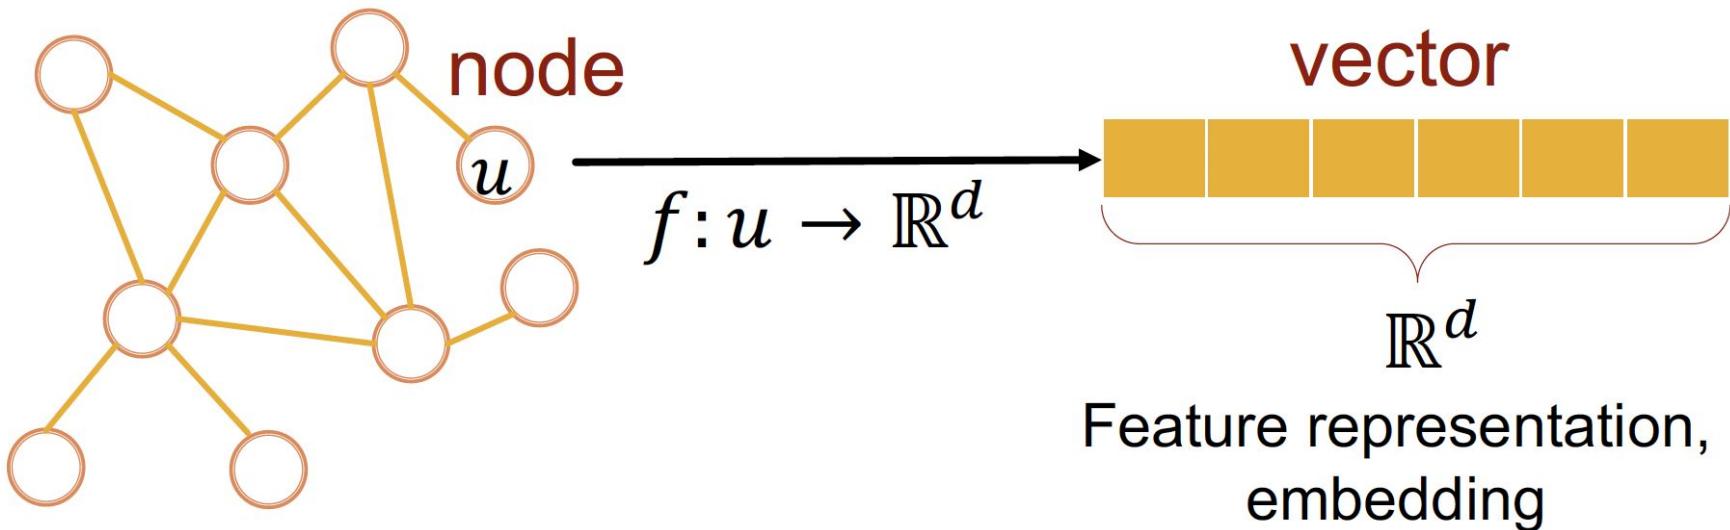
Graph representation learning

- Graph Representation Learning alleviates the need to do feature engineering **every single time.**



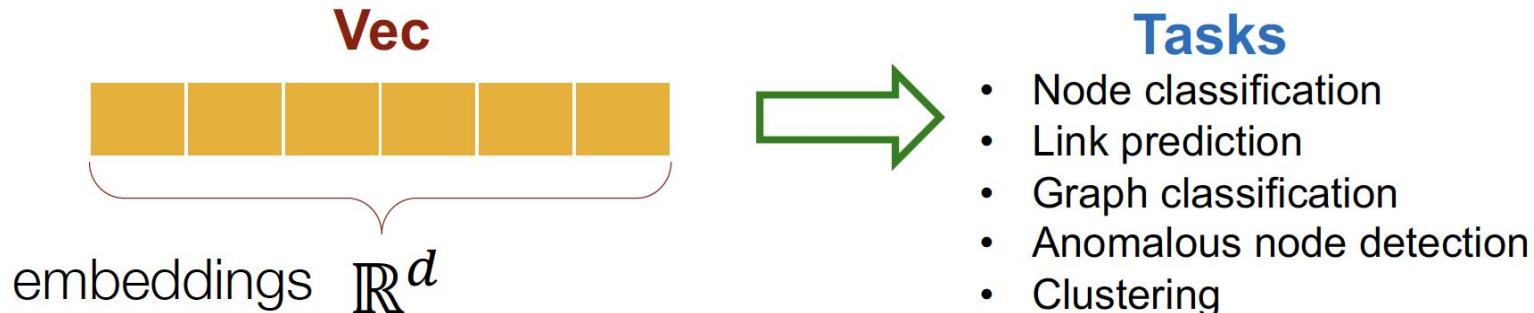
Graph representation learning

- Goal: Efficient task-independent feature learning for machine learning with graphs!



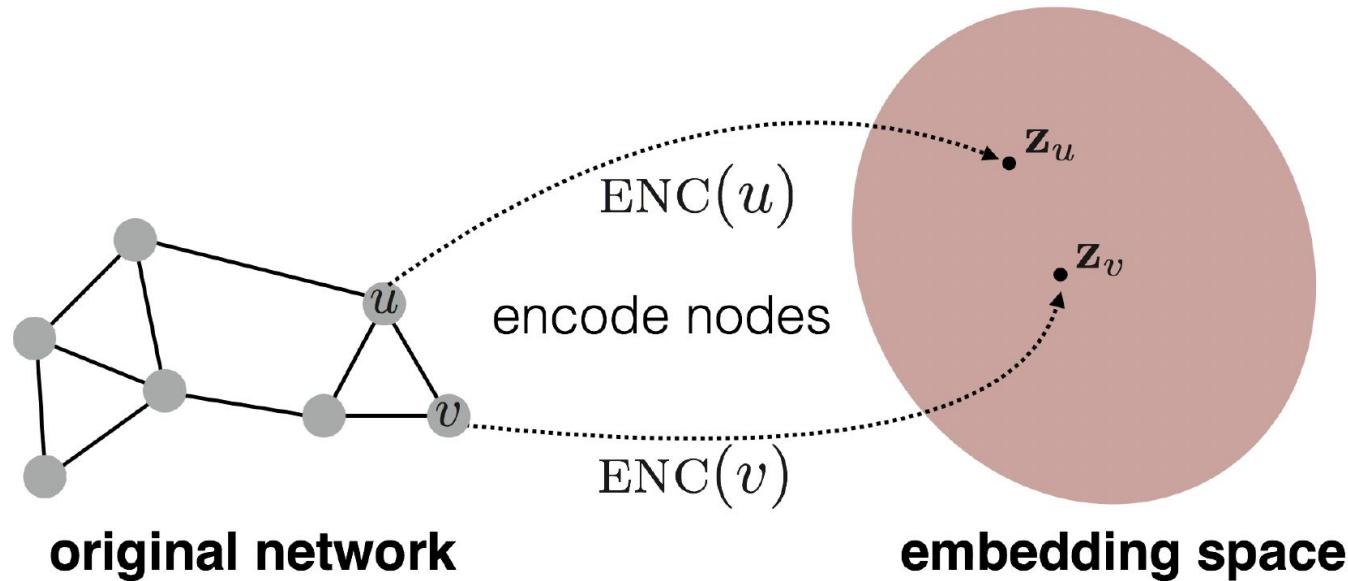
Why Embedding

- **Task:** Map nodes into an embedding space
- **Similarity of embeddings between nodes indicates**
 - their similarity in the network. For example:
 - Both nodes are close to each other (connected by an edge)
- **Encode network information**
- **Potentially used for many downstream predictions**



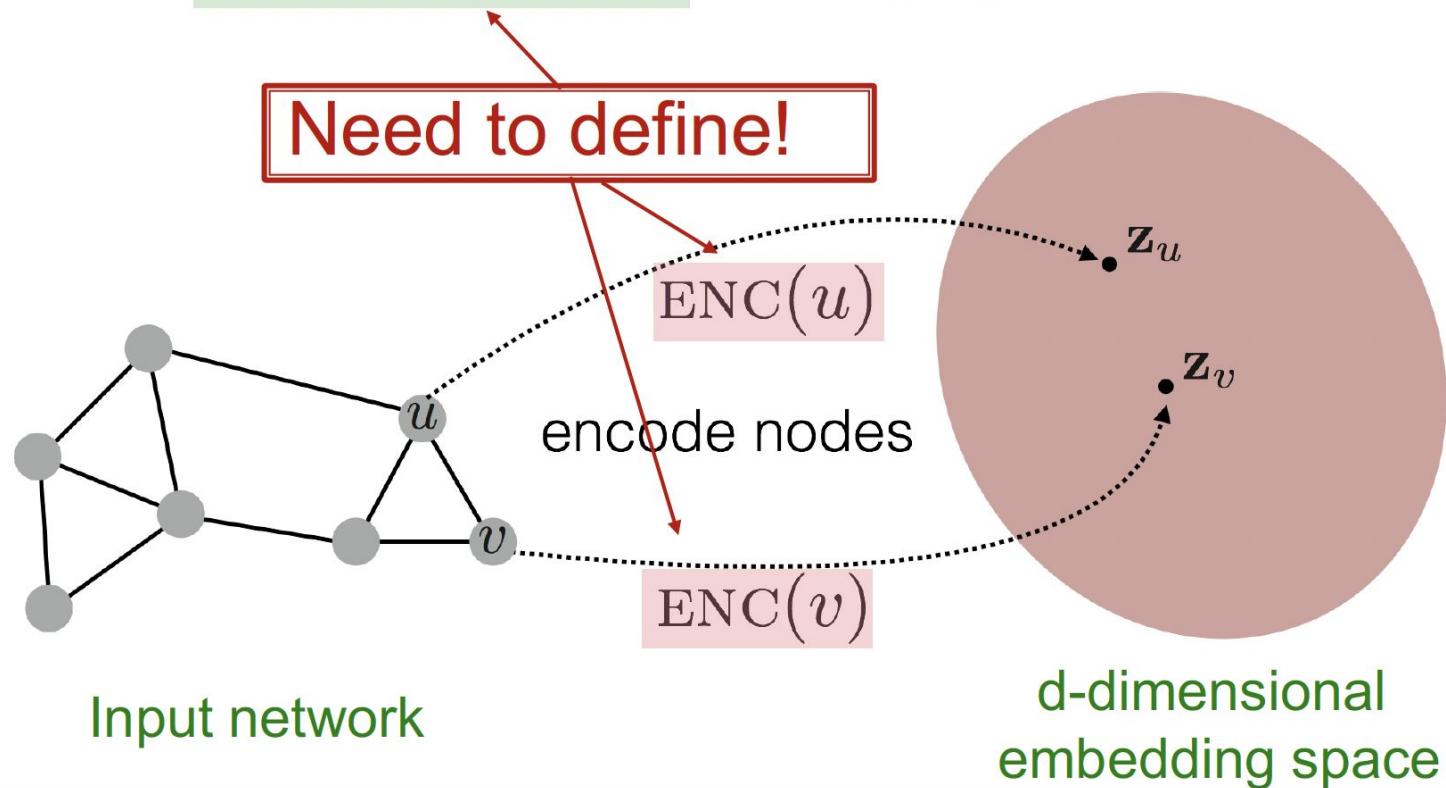
Embedding nodes

- Goal is to encode nodes so that **similarity in the embedding space** (e.g., dot product) approximates **similarity in the graph**



Embedding nodes

Goal: $\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$



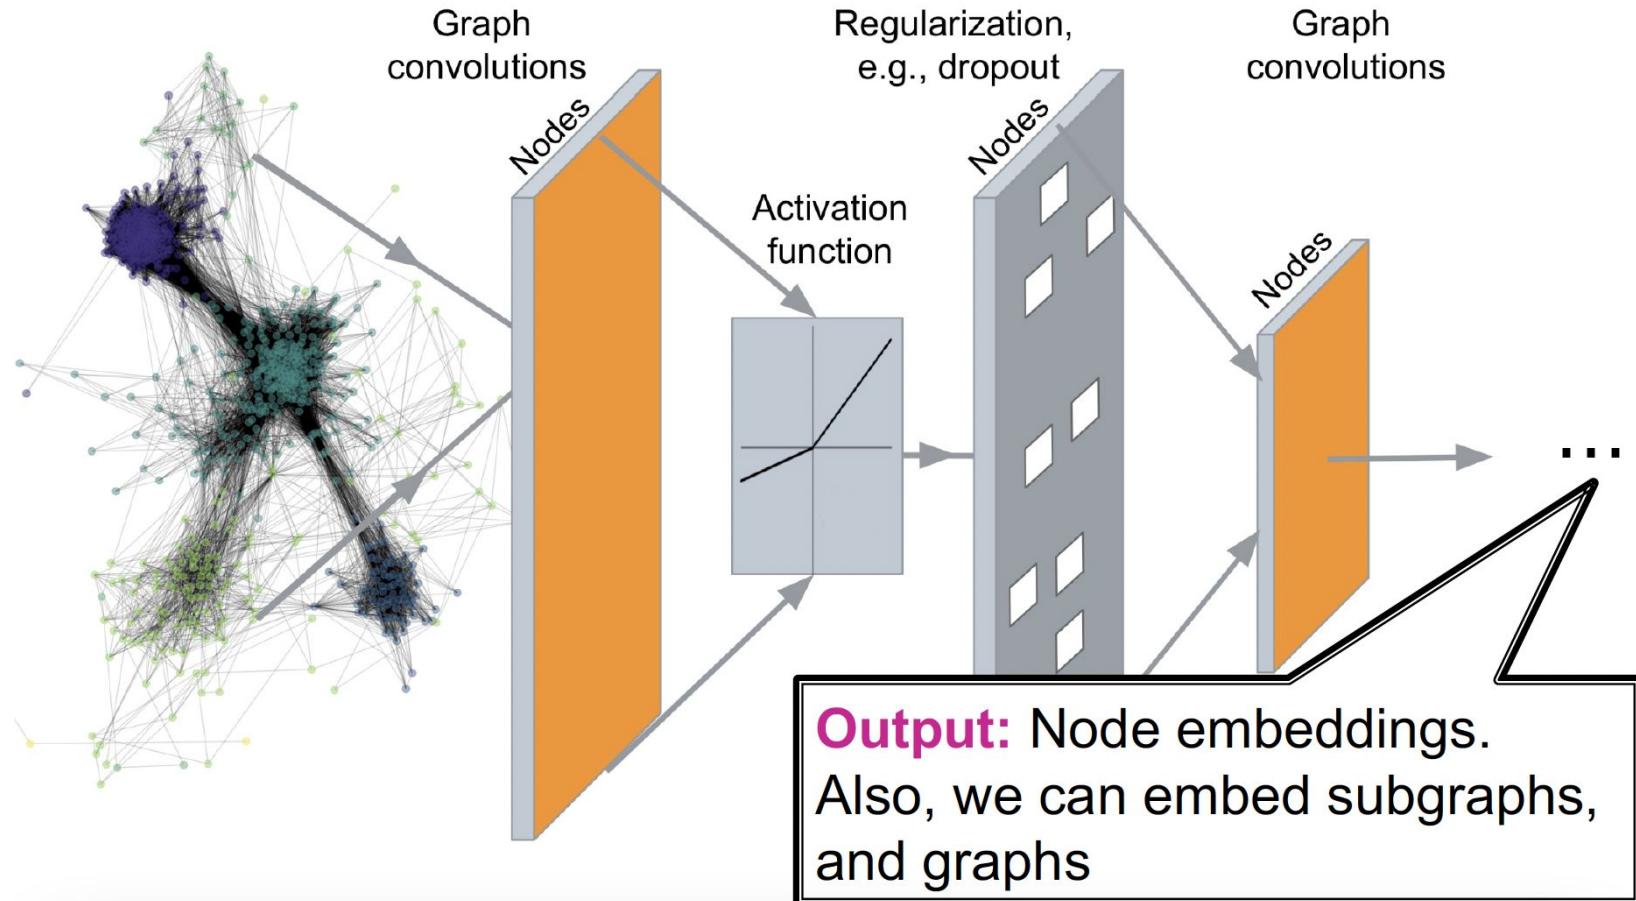
Deep graph encoders

- Deep learning methods based on **graph neural networks (GNNs)**:

$$\text{ENC}(v) =$$

**multiple layers of
non-linear transformations
based on graph structure**

Deep graph encoders



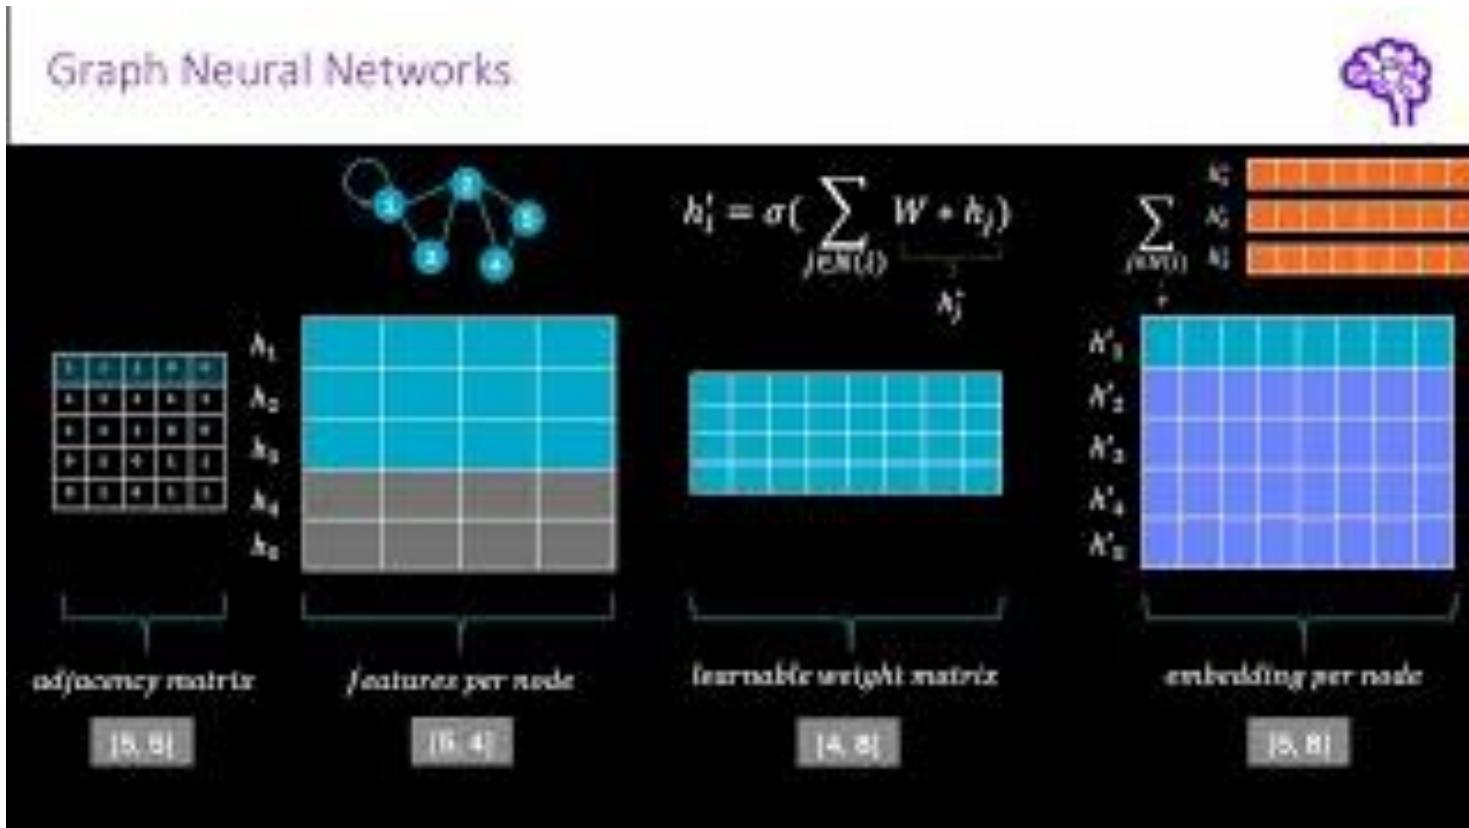
Recap: Basics of deep learning

- Loss function \mathcal{L} :

$$\min_{\Theta} \mathcal{L}(y, f_{\Theta}(x))$$

- f can be a simple linear layer, an MLP, or other neural networks (e.g., a GNN later)
- Sample a minibatch of input data x
- **Forward propagation:** Compute \mathcal{L} given x
- **Back-propagation:** Obtain gradient $\nabla_{\Theta} \mathcal{L}$ using a chain rule.
- Use **stochastic gradient descent (SGD)** to optimize \mathcal{L} for weights Θ over many iterations.

GNN functionality



<https://youtu.be/A-yKQamf2Fc?si=cfoIMTadg2UmDyyk>

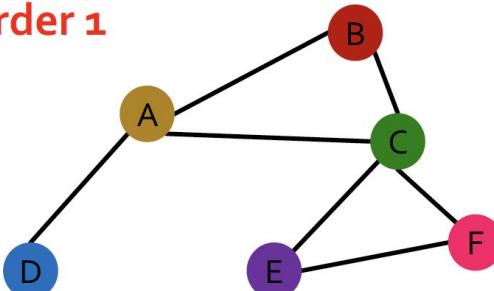
Permutation invariance

- Consider we want to embed an entire graph
- **Observation:** A graph does not have a canonical ordering of its nodes
 - We can have many different node orderings of the same graph.
- **What do we want:** If we learn an embedding function over a graph, we should get the same result (the same embedding) regardless of how the nodes are numbered.

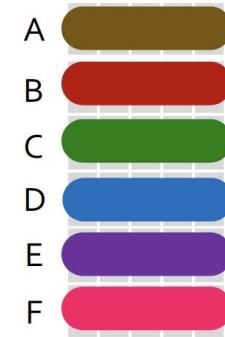
Permutation invariance

- Graph does not have a canonical ordering of the nodes!

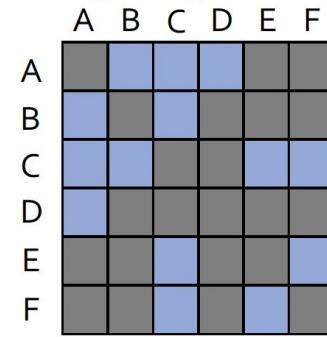
Order 1



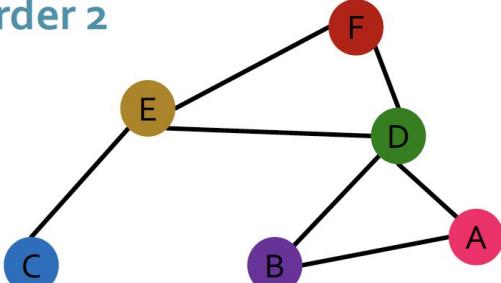
Node features X_1



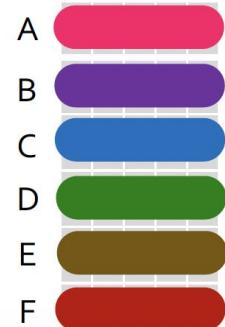
Adjacency matrix A_1



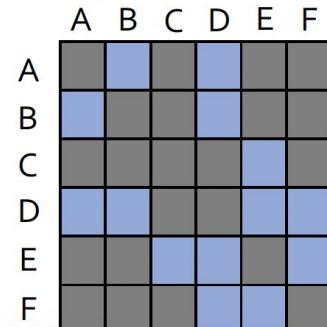
Order 2



Node features X_2



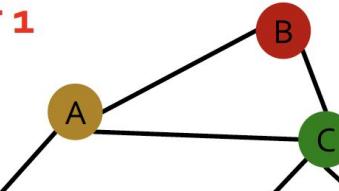
Adjacency matrix A_2



Permutation invariance

- Graph does not have a canonical ordering of the nodes!

Order 1



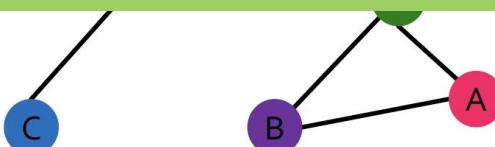
Node features X_1

A	[brown bar]
B	[red bar]
C	[green bar]
D	[blue bar]

Adjacency matrix A_1

	A	B	C	D	E	F
A	[grey]	[blue]	[blue]	[blue]	[grey]	[grey]
B	[blue]	[grey]	[grey]	[blue]	[grey]	[grey]
C	[blue]	[blue]	[grey]	[blue]	[blue]	[blue]

Learned graph representation
should be the same for Order 1
and Order 2



C	[blue bar]
D	[green bar]
E	[brown bar]
F	[red bar]

	C	D	E	F
C	[grey]	[blue]	[blue]	[blue]
D	[blue]	[grey]	[grey]	[blue]
E	[blue]	[blue]	[grey]	[blue]
F	[blue]	[blue]	[blue]	[grey]

2

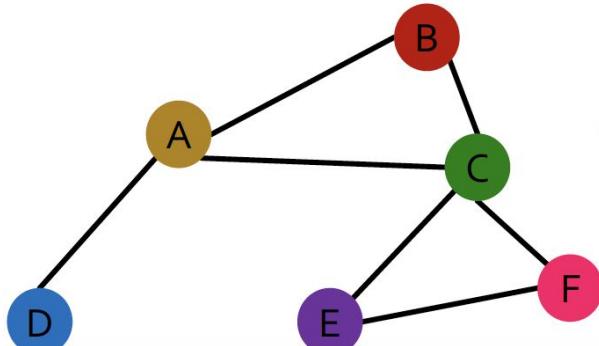
Permutation invariance

- What do we mean by “graph representation is the same for two orderings”?

$$f(\mathbf{A}_1, \mathbf{X}_1) = f(\mathbf{A}_2, \mathbf{X}_2)$$

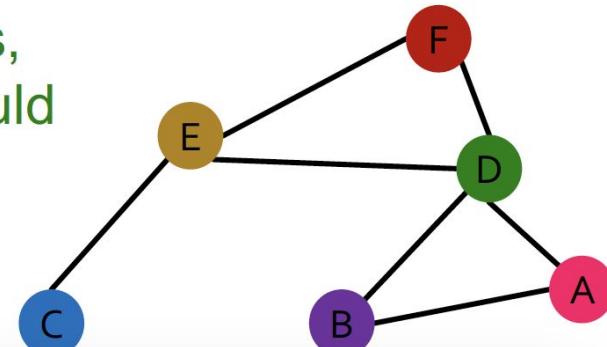
\mathbf{A} is the adjacency matrix
 \mathbf{X} is the node feature matrix

Order 1: $\mathbf{A}_1, \mathbf{X}_1$



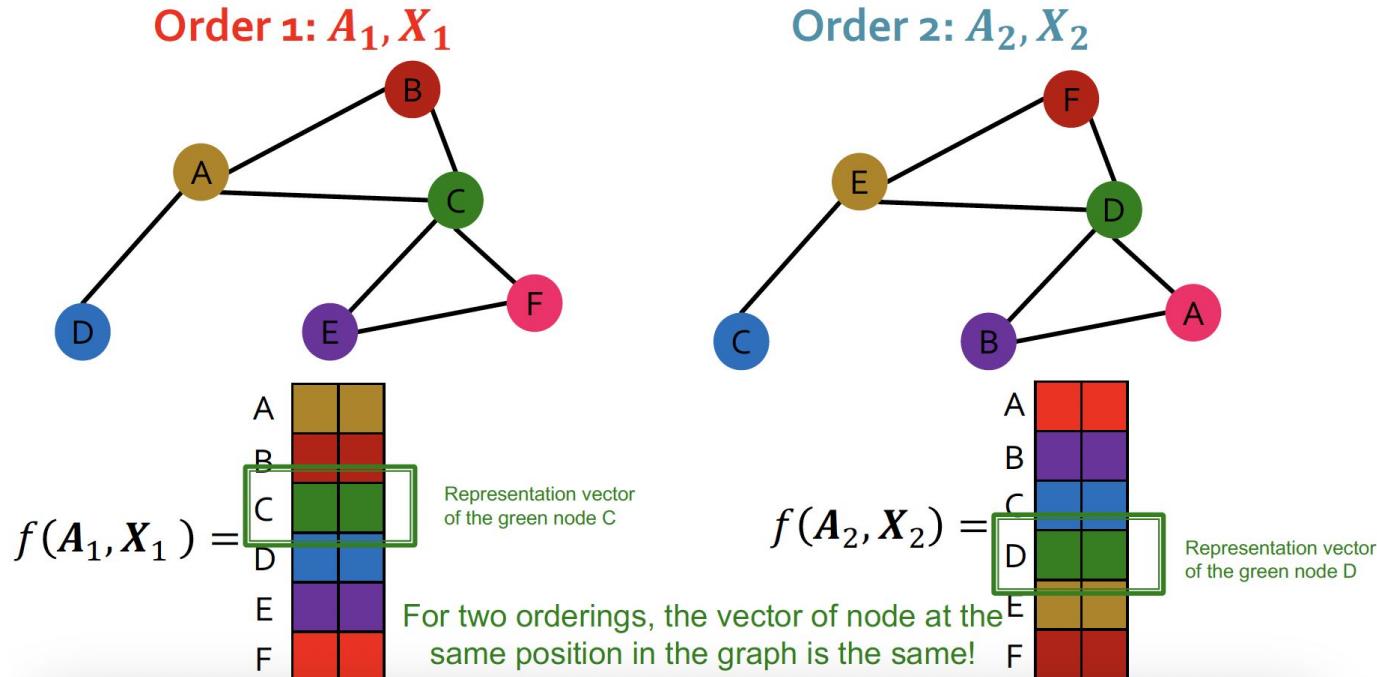
For two orders,
output of f should
be the same!

Order 2: $\mathbf{A}_2, \mathbf{X}_2$



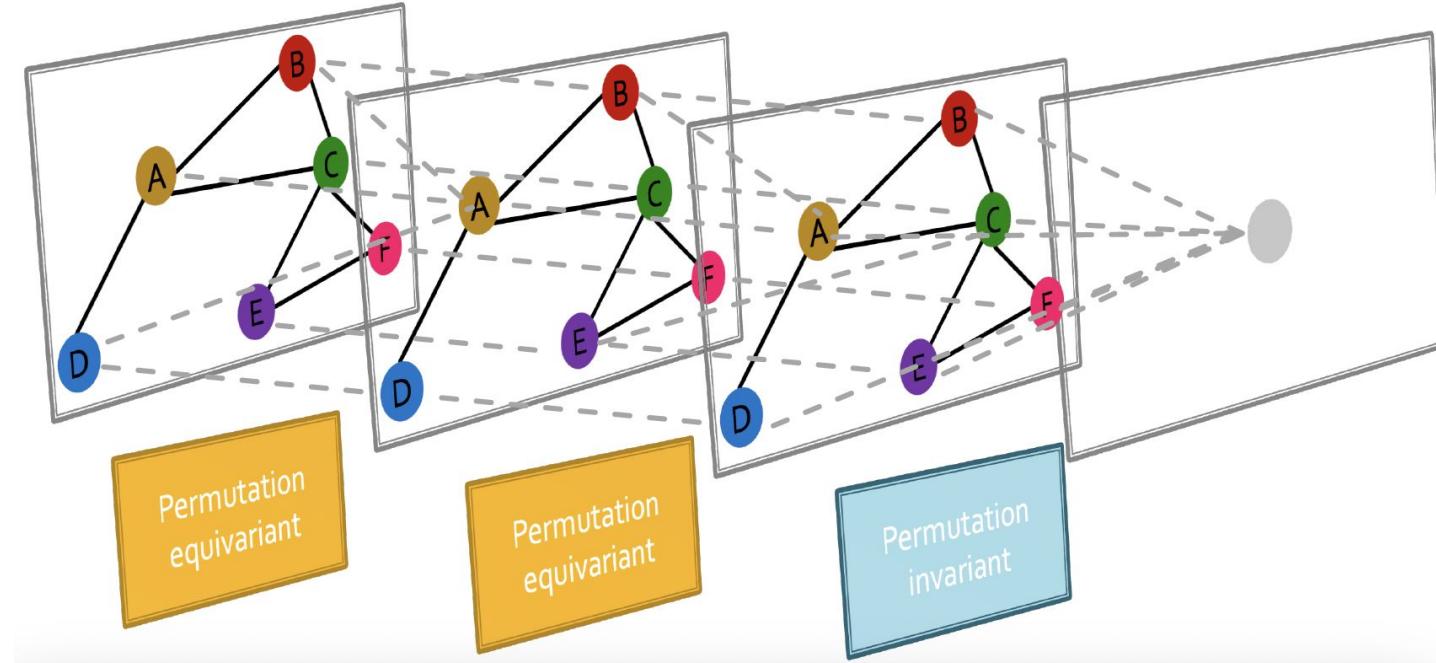
Permutation equivariance

- If the output vector of a node at the same position in the graph remains unchanged for any ordering



Graph neural networks overview

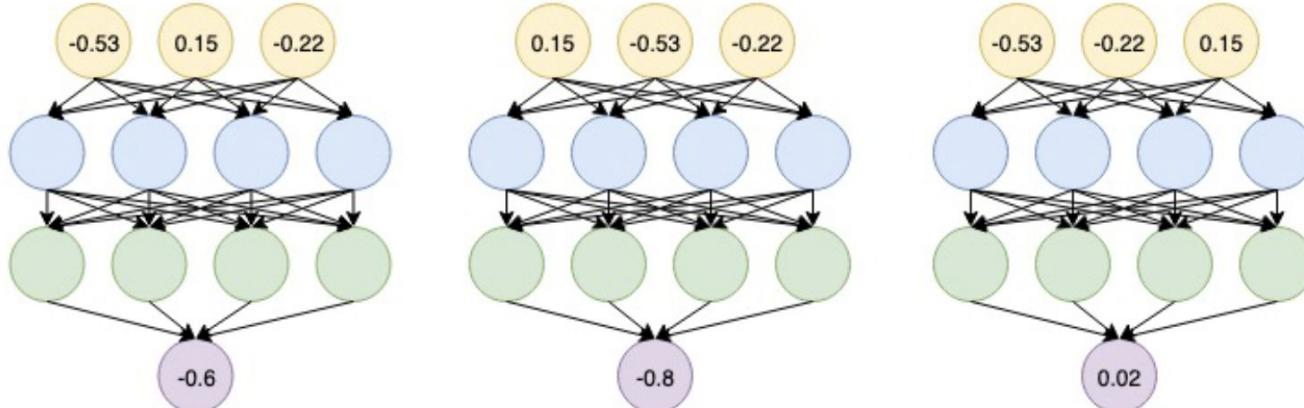
- Graph neural networks consist of multiple **permutation equivariant / invariant** functions



Graph neural networks overview

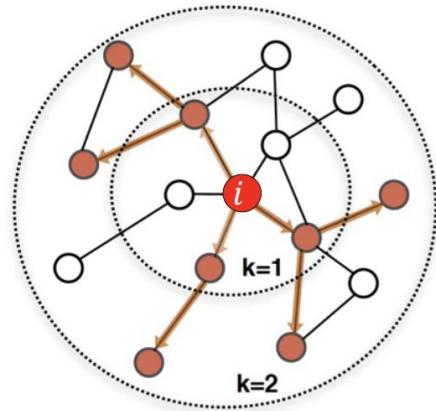
- Are other neural network architectures, e.g., MLPs, permutation invariant / equivariant? **NO!**

Switching the order of the input leads to different outputs!

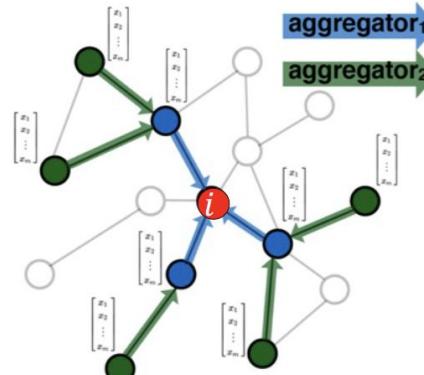


Graph convolutional networks

- Next: Design graph neural networks that are permutation invariant / equivariant by **passing and aggregating information from neighbors!**
- **IDEA:** Node's neighborhood defines a computation graph



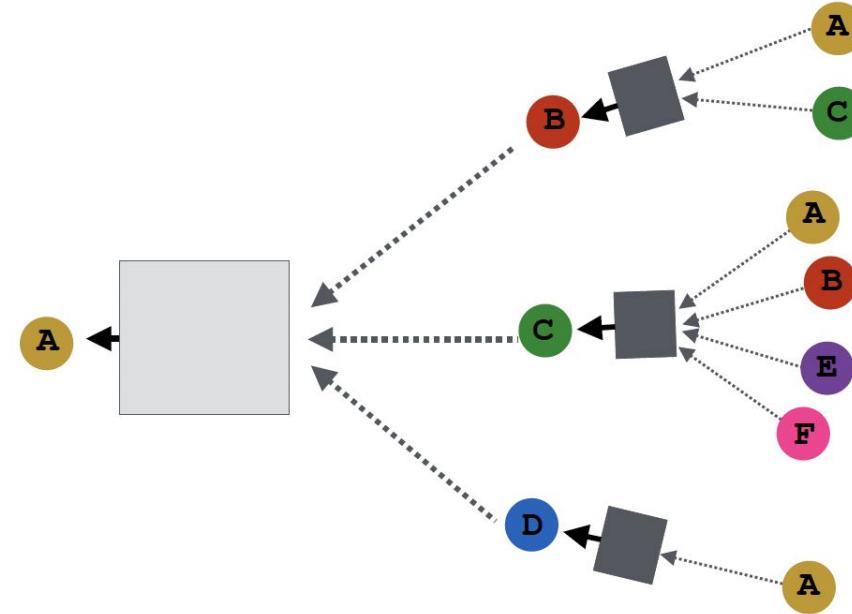
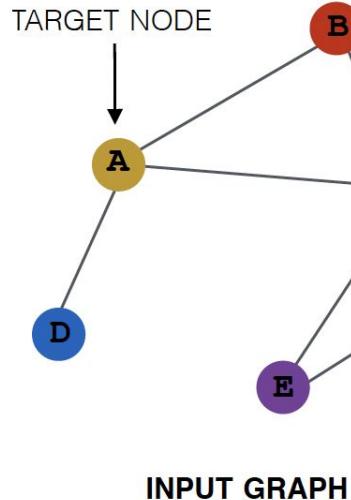
Determine node computation graph



Propagate and transform information

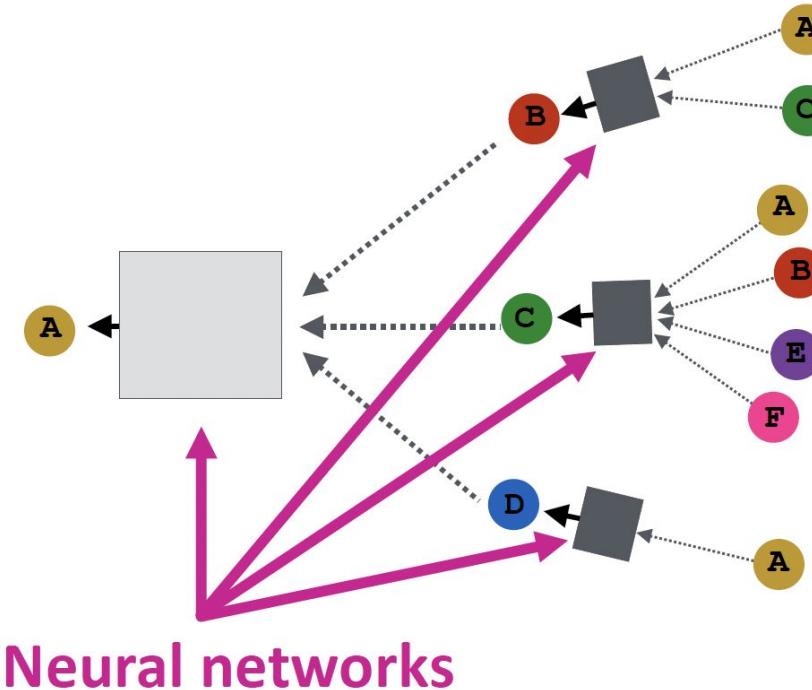
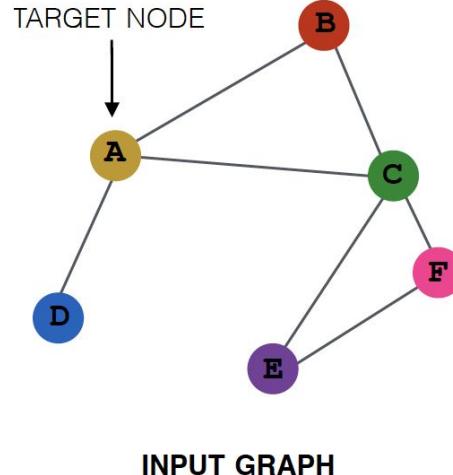
Aggregate neighbors

- Key idea: Generate node embeddings based on local network neighborhoods



Aggregate neighbors

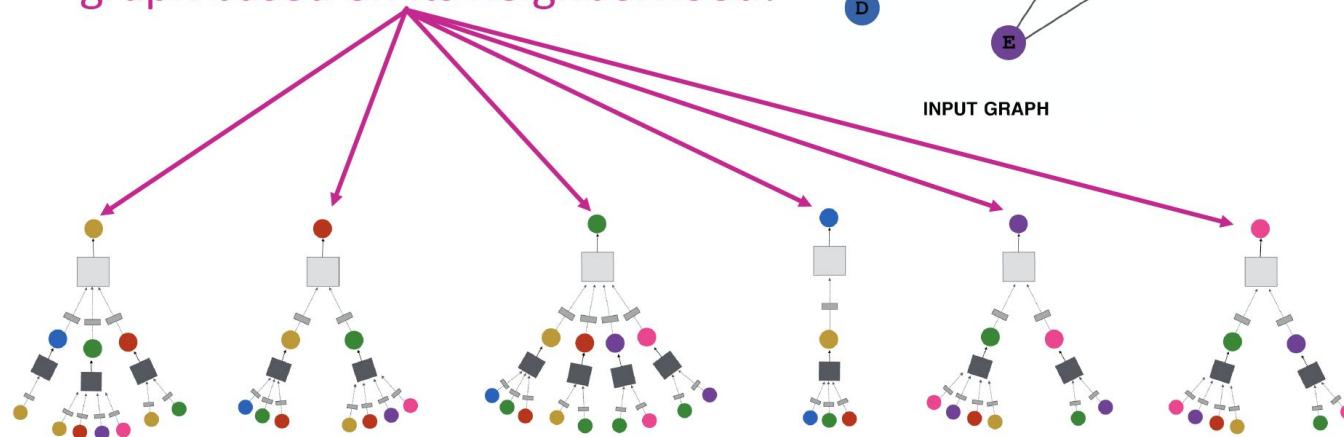
- **Intuition:** Nodes aggregate information from their neighbors using neural networks



Aggregate neighbors

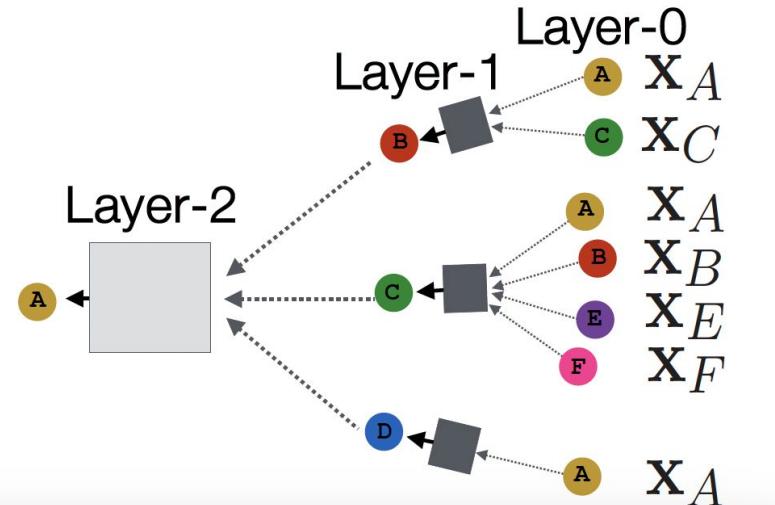
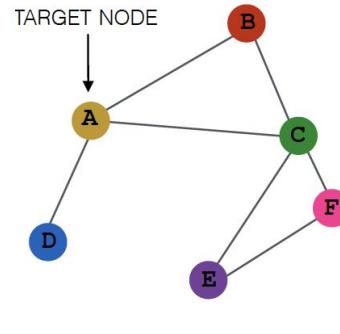
- **Intuition:** Network neighborhood defines a computation graph
- computation graph

Every node defines a computation graph based on its neighborhood!



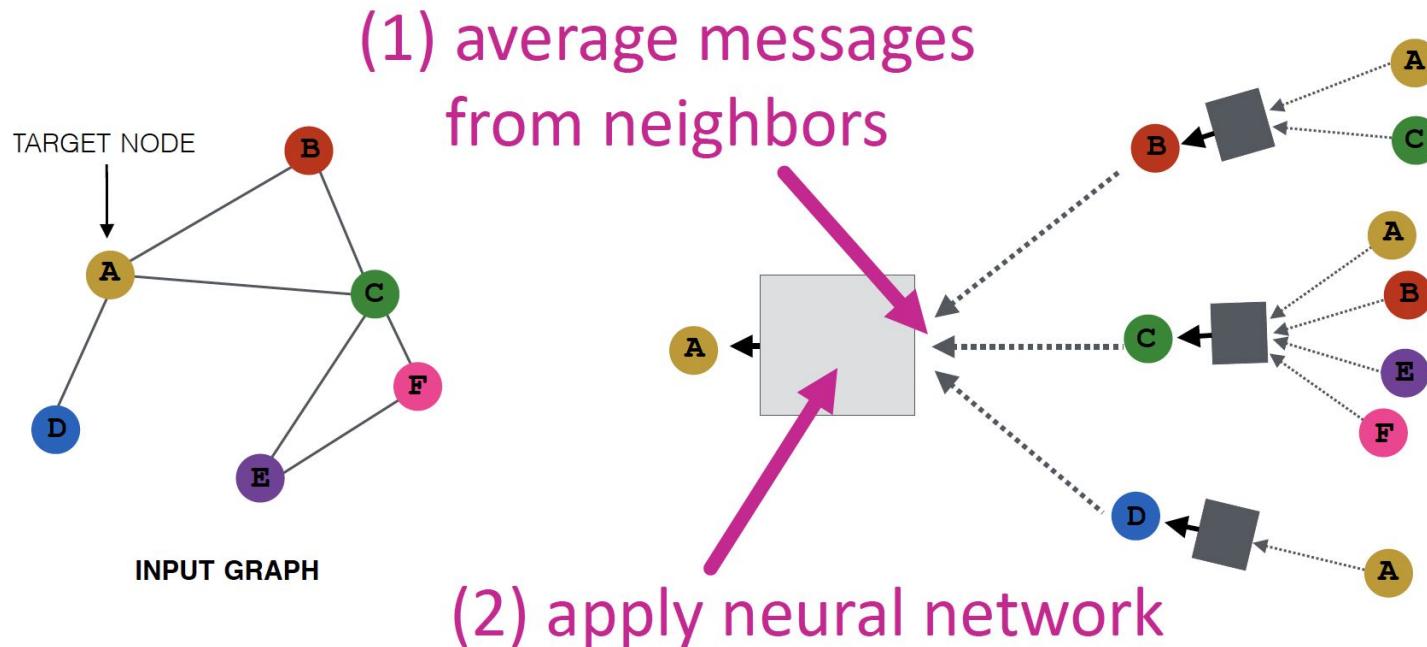
Deep model: many layers

- Model can be **of arbitrary depth**
 - Nodes have embeddings at each layer
 - Layer-0 embedding of node v is its input feature, x_v
 - Layer- k embedding gets information from nodes that are k hops away



Neighborhood aggregation

- **Basic approach:** Average information from neighbors and apply a neural network



The math: deep encoder of a GCN

- **Basic approach:** Average information from neighbors and apply a neural network

Initial 0-th layer embeddings are equal to node features

$$h_v^0 = x_v$$

embedding of v at layer k

$$h_v^{(k+1)} = \sigma(W_k \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} + B_k h_v^{(k)}), \forall k \in \{0, \dots, K-1\}$$

Total number of layers

Average of neighbor's previous layer embeddings

Non-linearity (e.g., ReLU)

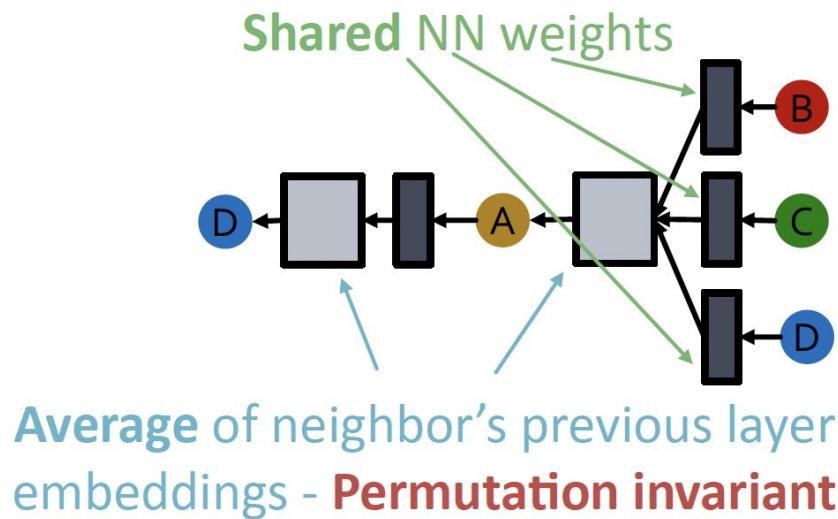
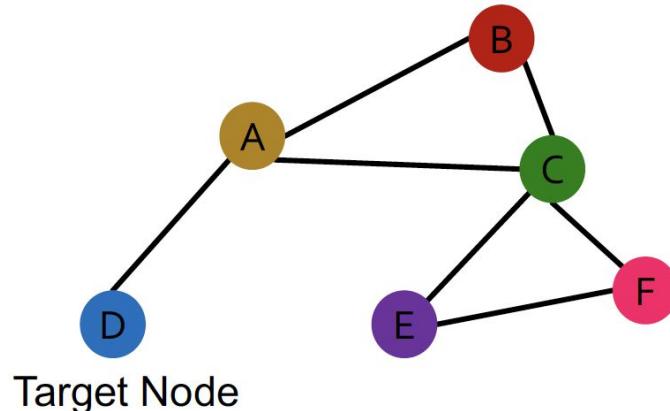
Notice summation is a permutation invariant pooling/aggregation.

Embedding after K layers of neighborhood aggregation

$$z_v = h_v^{(K)}$$

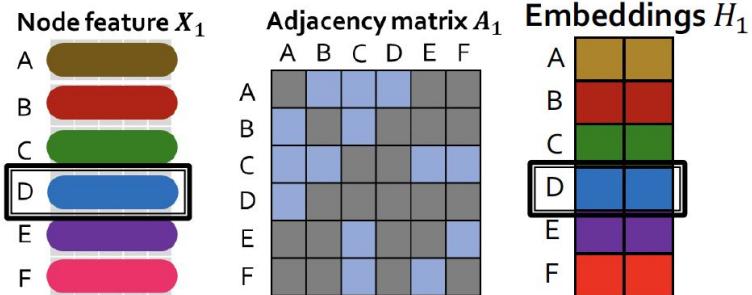
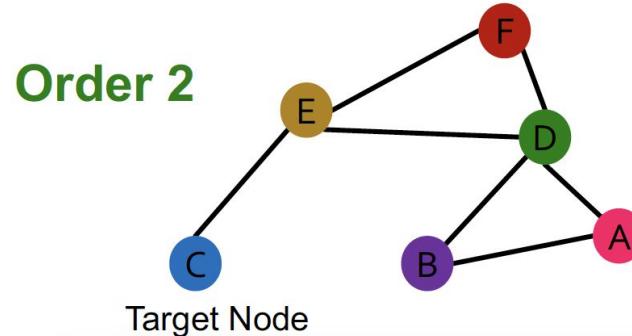
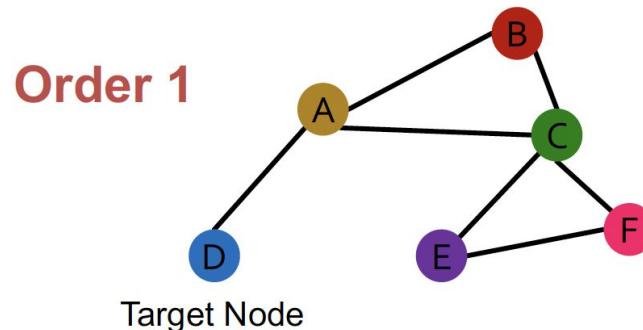
GCN: Invariance and equivariance

- What are the **invariance** and **equivariance** properties for a GCN?
- Given a node, the GCN that computes its embedding is **permutation invariant**

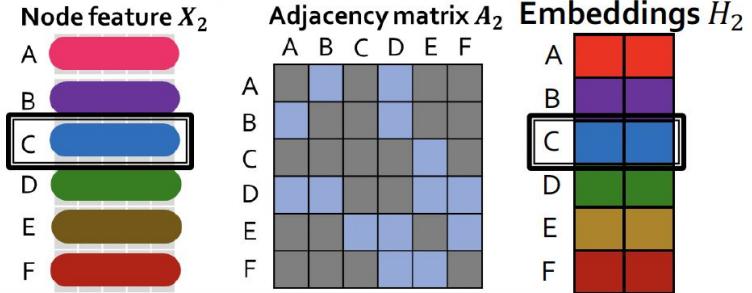


GCN: Invariance and equivariance

- Considering all nodes in a graph, GCN computation is **permutation equivariant!**



Permute the input, the output also permutes
accordingly - permutation equivariant



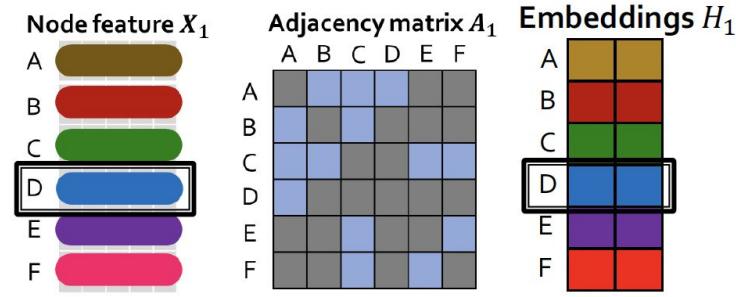
GCN: Invariance and equivariance

- Considering all nodes in a graph, GCN computation is **permutation equivariant!**

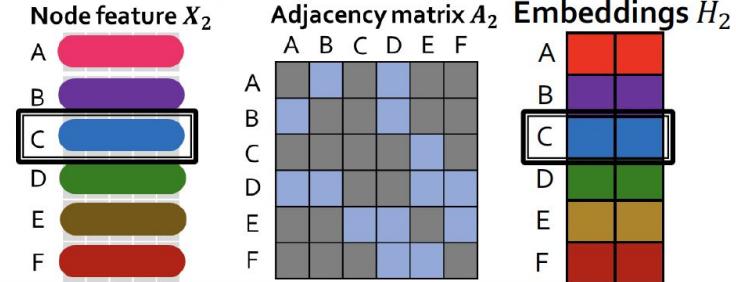
Detailed reasoning:

1. The rows of **input node features** and **output embeddings** are aligned
2. We know computing the embedding of a given node with GCN is **invariant**.
3. So, after permutation, the **location of a given node in the input node feature matrix** is changed, and the **the output embedding of a given node stays the same** (**the colors of node feature and embedding are matched**)

This is **permutation equivariant**



Permute the input, the output also permutes accordingly - **permutation equivariant**



Training the model

Trainable weight matrices
(i.e., what we learn)

$$\begin{aligned} h_v^{(0)} &= x_v \\ h_v^{(k+1)} &= \sigma(W_k \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} + B_k h_v^{(k)}), \forall k \in \{0..K-1\} \\ z_v &= h_v^{(K)} \end{aligned}$$

Final node embedding

We can feed these **embeddings into any loss function** and run SGD to **train the weight parameters**

h_v^k : the hidden representation of node v at layer k

- W_k : weight matrix for neighborhood aggregation
- B_k : weight matrix for transforming hidden vector of self

Training the model

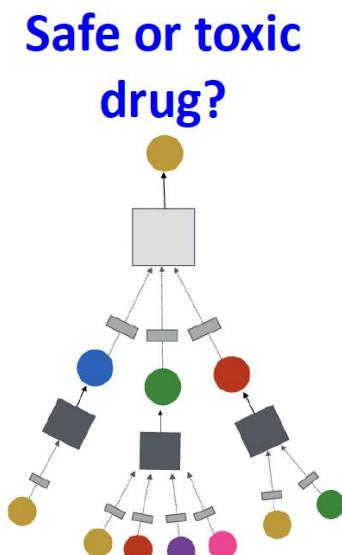
- Node embedding \mathbf{z}_v , is a function of input graph
- **Supervised setting:** We want to minimize loss \mathcal{L} :

$$\min_{\Theta} \mathcal{L}(y, f_{\Theta}(\mathbf{z}_v))$$

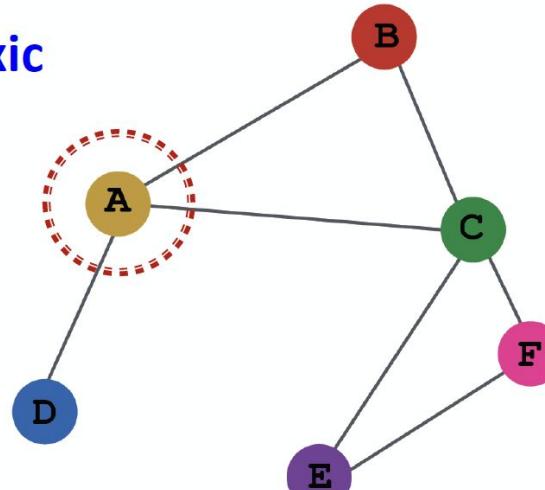
- y : node label
- \mathcal{L} could be L2 if y is real number, or cross entropy if y is categorical
- **Unsupervised setting:**
 - No node label available
 - **Use the graph structure as the supervision!**

Supervised training

- Directly train the model for a supervised task (e.g., node classification)



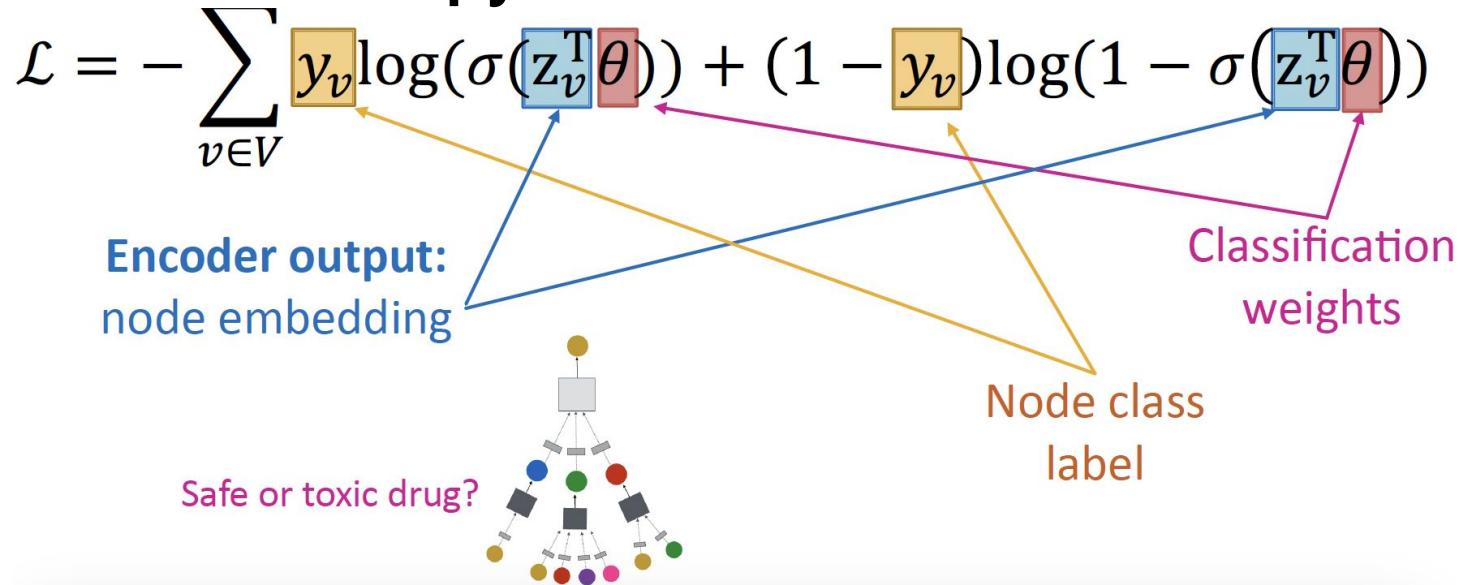
Safe or toxic drug?



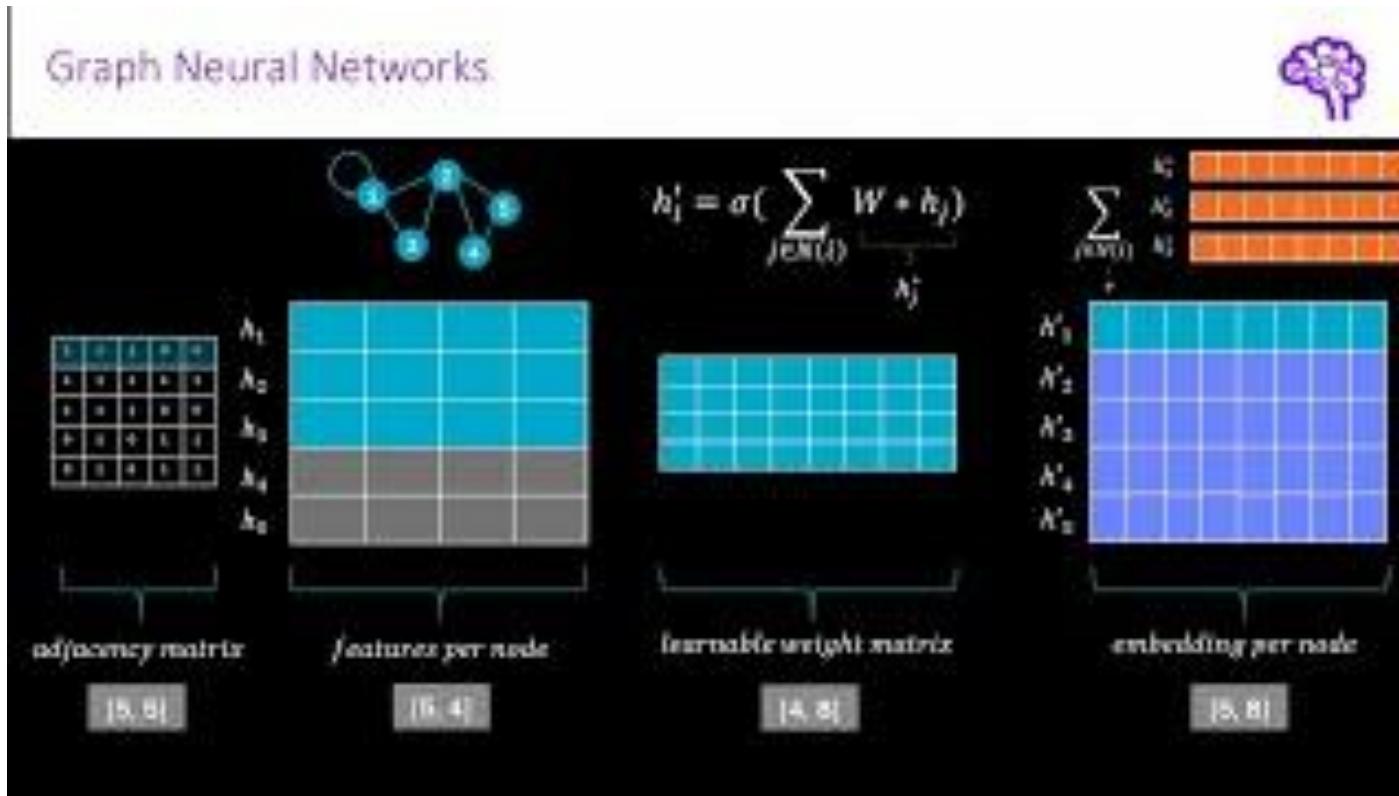
E.g., a drug-drug interaction network

Supervised training

- Directly train the model for a supervised task (e.g., node classification)
- Use cross entropy loss

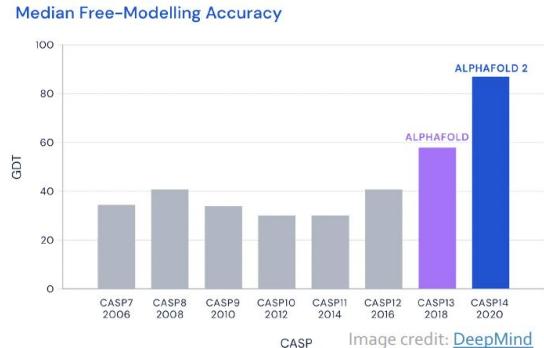


Graph Attention Networks!



<https://youtu.be/A-yKQamf2Fc?si=cfoIMTadg2UmDyyk>

Protein Folding



AlphaFold's AI could change the world of biological science as we know it

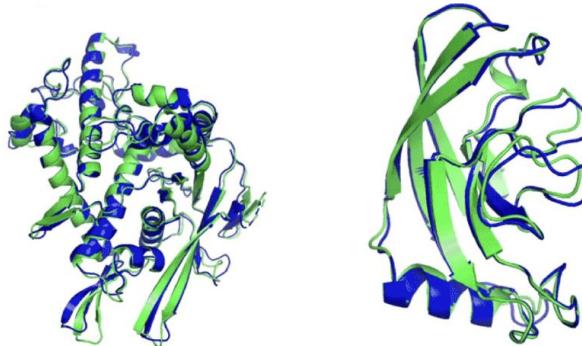
DeepMind's latest AI breakthrough can accurately predict the way proteins fold

Has Artificial Intelligence 'Solved' Biology's Protein-Folding Problem?

12-14-20
DeepMind's latest AI breakthrough could turbocharge drug discovery

Protein Folding

- Computational predict a protein's **3D structure** based solely on its **amino acid sequence**
 - For each graph node predict its 3D coordinates



T1037 / 6vr4
90.7 GDT
(RNA polymerase domain)

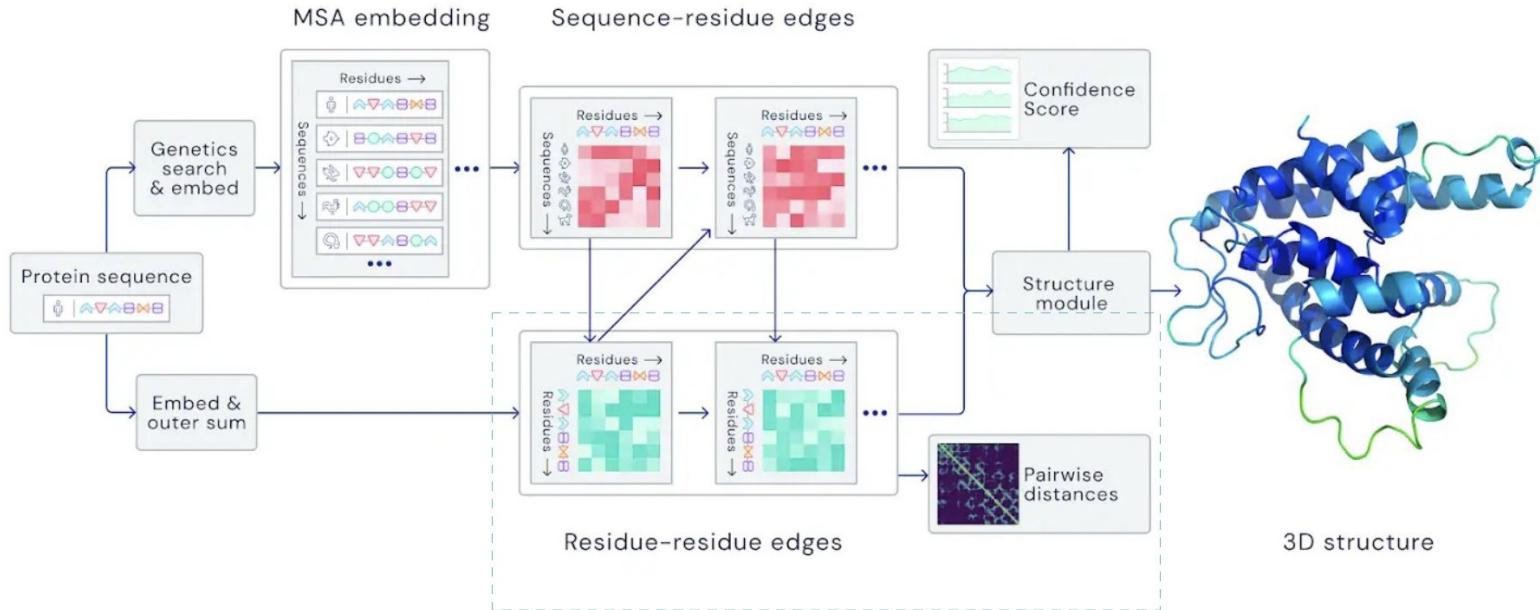
T1049 / 6y4f
93.3 GDT
(adhesin tip)

- Experimental result
- Computational prediction

Image credit: [DeepMind](#)

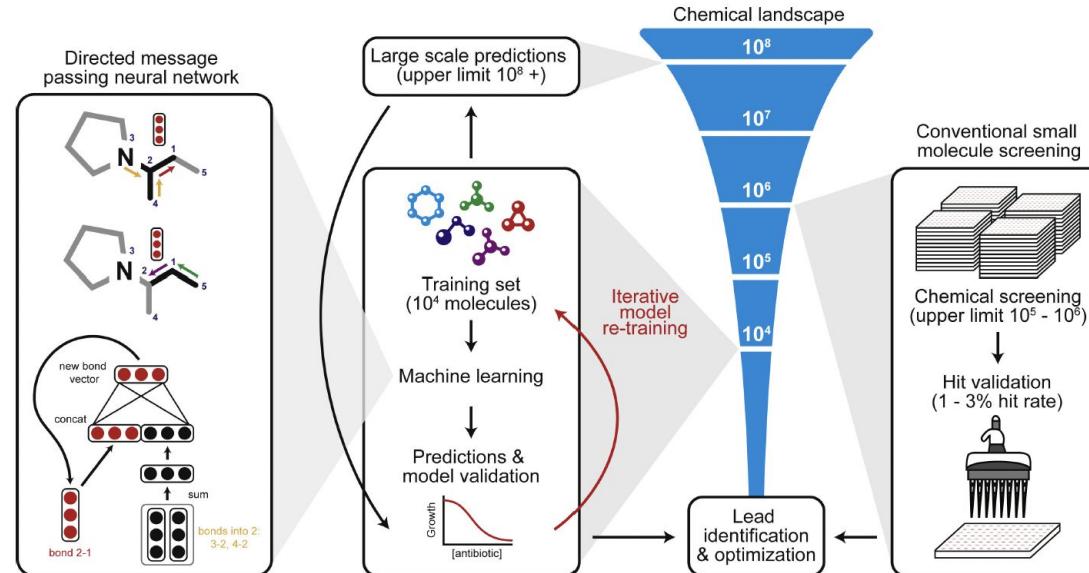
AlphaFold: Solving Protein Folding

- **Nodes:** Amino acids in a protein sequence
- **Edges:** Proximity between amino acids (residues)



GNN for Antibiotic Discovery

- A GNN **graph classification model**
- Predict promising molecules (**hits**) from a pool of candidates



Graph machine learning tools

- **PyG (PyTorch Geometric)**
 - The ultimate library for Graph Neural Networks (GNN).
- **GraphGym**
 - Platform for designing graph neural networks.
 - Modularized GNN implementation, simple hyperparameter tuning, flexible user customization